




arXiv is now an independent nonprofit! [Learn more](#)

arxiv





[Search](#) [Submit](#) [Donate](#) [Log in](#)

Showing 201–301 of 301 results

Search v0.5.6 released 2020-02-24

Query: order: -announced_date_first; size: 200; page_start: 200;
date_range: from 2019-01-01 to 2026-09-20; include_cross_list: True;
terms: AND all="zero-shot" OR "few-shot" OR "low-resource" OR
"limited labelled data" OR "limited labeled data" OR "large language
model*" OR "prompt-based" OR verbalizer OR "fine-tun*" OR "task-
specific" OR "natural language inference"; AND all="text
classification" OR "stance detection" OR "sentiment classification"
OR "sentiment analysis" OR "content analysis" OR annotation; AND
all="social media" OR reddit OR twitter OR "user-generated content"

[Simple Search](#)

OR "online discourse" OR forum*; AND all=compar* OR benchmark* OR evaluat*

Refine query

New search

200 ▾

results per page. Sort results by

Announcement date (newest first) ▾

Go

Previous

1

2

201. [arXiv:2407.00167](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.ET](#) [cs.HC](#) [cs.SI](#)

Can GPT-4 Help Detect Quit Vaping Intentions? An Exploration of Automatic Data **Annotation** Approach

Authors: [Sai Krishna Revanth Vuruma](#), [Dezhi Wu](#), [Saborny Sen Gupta](#), [Lucas Aust](#), [Valerie Lookingbill](#), [Wyatt Bellamy](#), [Yang Ren](#), [Erin Kasson](#), [Li-Shiun Chen](#), [Patricia Cavazos-Rehg](#), [Dian Hu](#), [Ming Huang](#)

Abstract: ...and fatalities during the EVALI outbreak in 2019, highlighting the urgency to comprehend vaping behaviors and develop effective strategies for cessation. Due to the ubiquity of **social**... ▾ [More](#)

Submitted 28 June, 2024; **originally announced** July 2024.

Comments: Accepted for the AI Applications in Public Health and Social Services workshop at the 22nd International Conference on Artificial Intelligence in Medicine (AIME 2024)

202. [arXiv:2407.00132](#) [[pdf](#), [other](#)] [cs.SE](#) [cs.AI](#)

ShortcutsBench: A Large-Scale Real-world **Benchmark** for API-based Agents

Authors: [Haiyang Shen](#), [Yue Li](#), [Desong Meng](#), [Dongqi Cai](#), [Sheng Qi](#), [Li Zhang](#), [Mengwei Xu](#), [Yun Ma](#)

Abstract: Recent advancements in integrating **large**... ▾ [More](#)

Submitted 23 January, 2025; **v1 submitted** 28 June, 2024; **originally announced** July 2024.

Comments: ICLR'25: <https://openreview.net/forum?id=kKILfPkhSz>

203. [arXiv:2406.18239](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Zero-shot prompt-based classification: topic labeling in times of foundation models in German Tweets

Authors: [Simon Munker](#), [Kai Kugler](#), [Achim Rettinger](#)

Abstract: Filtering and **annotating** textual data are routine tasks in many areas, like... ▾ [More](#)

Submitted 26 June, 2024; **originally announced** June 2024.

Comments: 10 pages, 2 tables, 1 figure

204. [arXiv:2406.07594](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CR](#)

MLLMGuard: A Multi-dimensional Safety Evaluation Suite for Multimodal **Large Language Models**

Authors: [Tianle Gu](#), [Zeyang Zhou](#), [Kexin Huang](#), [Dandan Liang](#), [Yixu Wang](#), [Haiquan Zhao](#), [Yuanqi Yao](#), [Xingge Qiao](#), [Keqing Wang](#), [Yujiu Yang](#), [Yan Teng](#), [Yu Qiao](#), [Yingchun Wang](#)

Abstract: Powered by remarkable advancements in **Large**... ▾ [More](#)

Submitted 13 June, 2024; **v1 submitted** 11 June, 2024; **originally announced** June 2024.

205. [arXiv:2406.07483](#) [[pdf](#), [other](#)] [cs.CL](#)

Advancing Annotation of Stance in Social Media Posts: A Comparative Analysis of Large Language Models and Crowd Sourcing

Authors: [Mao Li](#), [Frederick Conrad](#)

Abstract: In the rapidly evolving landscape of Natural Language Processing (NLP), the use of **Large...** [▽ More](#)

Submitted 11 June, 2024; **originally announced** June 2024.

206. [arXiv:2406.04892](#) [[pdf](#), [other](#)] [cs.CL](#)

Sexism Detection on a Data Diet

Authors: [Rabiraj Bandyopadhyay](#), [Dennis Assenmacher](#), [Jose M. Alonso Moral](#), [Claudia Wagner](#)

Abstract: There is an increase in the proliferation of online hate commensurate with the rise in the usage of **social...** [▽ More](#)

Submitted 7 June, 2024; **originally announced** June 2024.

Comments: Accepted at ACM WebSci 2024 Workshop in DHOW: Diffusion of Harmful Content on Online Web Workshop

207. [arXiv:2405.15145](#) [[pdf](#), [other](#)] [cs.AI](#) [cs.CL](#) [cs.MA](#)

CulturePark: Boosting Cross-cultural Understanding in Large Language Models

Authors: [Cheng Li](#), [Damien Teney](#), [Linyi Yang](#), [Qingsong Wen](#), [Xing Xie](#), [Jindong Wang](#)

Abstract: Cultural bias is pervasive in many **large...** [▽ More](#)

Submitted 21 November, 2024; **v1** submitted 23 May, 2024; **originally announced** May 2024.

Comments: NeurIPS 2024; Code is released at <https://github.com/Scarelette/CulturePark>. arXiv admin note: substantial text overlap with arXiv:2402.10946

208. [arXiv:2405.13056](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.SI](#)

Large language models for newspaper sentiment analysis during COVID-19: The Guardian

Authors: [Rohitash Chandra](#), [Baicheng Zhu](#), [Qingying Fang](#), [Eka Shinjikashvili](#)

Abstract: ...coverage encompassed a wide range of topics that includes viral transmission, allocation of medical resources, and government response measures. There have been studies on **sentiment...** [▽ More](#)

Submitted 27 April, 2025; **v1** submitted 20 May, 2024; **originally announced** May 2024.

Journal ref: Applied Soft Computing, Volume 171, 2025, pp. 112743

209. [arXiv:2405.12929](#) [[pdf](#), [other](#)] [cs.CL](#) [doi](#) [10.1109/TAFFC.2025.3553399](#)

Code-mixed Sentiment and Hate-speech Prediction

Authors: [Anjali Yadav](#), [Tanya Garg](#), [Matej Klemen](#), [Matej Ulcar](#), [Basant Agarwal](#), [Marko Robnik Sikonja](#)

Abstract: ...used in informal discourse in countries with several official languages, but also in many other countries in combination with English or neighboring languages. As recently **large...** [▽ More](#)

Submitted 21 May, 2024; **originally announced** May 2024.

210. [arXiv:2405.10700](#) [[pdf](#), [other](#)] [cs.IR](#) [cs.AI](#) [cs.CL](#) [cs.CY](#) [doi](#) [10.1145/3626772.3657667](#)

SynDy: Synthetic Dynamic Dataset Generation Framework for Misinformation Tasks

Authors: [Michael Shlisselberg](#), [Ashkan Kazemi](#), [Scott A. Hale](#), [Shiri Dori-Hacohen](#)

Abstract: ...of nascent fact-checking initiatives. In this paper we present SynDy, a framework for Synthetic Dynamic Dataset Generation to leverage the capabilities of the largest frontier **Large...**

[▽ More](#)

Submitted 17 May, 2024; **originally announced** May 2024.

Journal ref: Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '24), July 14--18, 2024, Washington, DC, USA

211. [arXiv:2405.06239](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

SaudiBERT: A **Large Language Model** Pretrained on Saudi Dialect Corpora

Authors: [Faisal Qarah](#)

Abstract: In this paper, we introduce SaudiBERT, a monodialect Arabic language model pretrained exclusively on Saudi dialectal text. To demonstrate the model's effectiveness, we **compared** SaudiBERT with six different multidialect Arabic language models across 11 evaluation datasets, which are divided into two groups:... [▽ More](#)

Submitted 10 May, 2024; **originally announced** May 2024.

212. [arXiv:2405.03688](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#)

Large Language Models Reveal Information Operation Goals, Tactics, and Narrative Frames

Authors: [Keith Burghardt](#), [Kai Chen](#), [Kristina Lerman](#)

Abstract: ...is limited by manual analysis of messages and subjective interpretation of their observable behavior. In this paper, we explore whether these limitations can be mitigated with **large...** [▽ More](#)

Submitted 6 May, 2024; **originally announced** May 2024.

Comments: 15 pages, 9 figures

213. [arXiv:2404.18371](#) [[pdf](#), [other](#)] [cs.CL](#)

QANA: LLM-based Question Generation and Network Analysis for **Zero-shot** Key Point Analysis and Beyond

Authors: [Tomoki Fukuma](#), [Koki Noda](#), [Toshihide Ubukata](#) [Kousuke Hoso](#), [Yoshiharu Ichikawa](#), [Kyosuke Kambe](#), [Yu Masubuchi](#), [Fujio Toriumi](#)

Abstract: The proliferation of **social...** [▽ More](#)

Submitted 28 April, 2024; **originally announced** April 2024.

Comments: Under review as a conference paper at COLM 2024

214. [arXiv:2404.17607](#) [[pdf](#), [other](#)] [cs.IR](#) [cs.AI](#) [cs.CL](#) [cs.LG](#) [cs.SI](#)

Utilizing **Large Language Models** to Identify **Reddit** Users Considering Vaping Cessation for Digital Interventions

Authors: [Sai Krishna Revanth Vuruma](#), [Dezhi Wu](#), [Saborny Sen Gupta](#), [Lucas Aust](#), [Valerie Lookingbill](#), [Caleb Henry](#), [Yang Ren](#), [Erin Kasson](#), [Li-Shiun Chen](#), [Patricia Cavazos-Rehg](#), [Dian Hu](#), [Ming Huang](#)

Abstract: The widespread adoption of **social...** [▽ More](#)

Submitted 25 April, 2024; **originally announced** April 2024.

215. [arXiv:2404.16461](#) [pdf, other] cs.CL

Large Language Models Perform on Par with Experts Identifying Mental Health Factors in Adolescent Online Forums

Authors: [Isabelle Lorge](#), [Dan W. Joyce](#), [Andrey Kormilitzin](#)

Abstract: Mental health in children and adolescents has been steadily deteriorating over the past few years. The recent advent of **Large...** [▽ More](#)

Submitted 26 April, 2024; v1 submitted 25 April, 2024; originally announced April 2024.

216. [arXiv:2404.16116](#) [pdf, other] cs.CL cs.AI

Classifying Human-Generated and AI-Generated Election Claims in Social Media

Authors: [Alphaeus Dmonte](#), [Marcos Zampieri](#), [Kevin Lybarger](#), [Massimiliano Albanese](#), [Genya Coulter](#)

Abstract: Politics is one of the most prevalent topics discussed on **social...** [▽ More](#)

Submitted 25 April, 2024; v1 submitted 24 April, 2024; originally announced April 2024.

217. [arXiv:2404.14339](#) [pdf, other] cs.CL

Zero-shot Cross-lingual Stance Detection via Adversarial Language Adaptation

Authors: [Bharathi A](#), [Arkaitz Zubiaga](#)

Abstract: **Stance...** [▽ More](#)

Submitted 22 April, 2024; originally announced April 2024.

218. [arXiv:2404.12171](#) [pdf, other] cs.CL cs.SI

Stance Detection on Social Media with Fine-Tuned Large Language Models

Authors: [İlker Gül](#), [Rémi Lebret](#), [Karl Aberer](#)

Abstract: **Stance...** [▽ More](#)

Submitted 18 April, 2024; originally announced April 2024.

219. [arXiv:2404.01753](#) [pdf, other] cs.CL

M2SA: Multimodal and Multilingual Model for Sentiment Analysis of Tweets

Authors: [Gaurish Thakkar](#), [Sherzod Hakimov](#), [Marko Tadić](#)

Abstract: ...has garnered significant attention. However, there needs to be more clarity when it comes to analysing multimodal tasks in multi-lingual contexts. While prior studies on **sentiment...** [▽ More](#)

Submitted 12 June, 2024; v1 submitted 2 April, 2024; originally announced April 2024.

Journal ref: LREC-COLING 2024 - The 2024 Joint International Conference on Computational Linguistics, Language Resources and **Evaluation**

220. [arXiv:2403.19031](#) [pdf] cs.CL cs.AI cs.LG

Evaluating Large Language Models for Health-Related Text Classification Tasks with Public Social Media Data

Authors: [Yuting Guo](#), [Anthony Ovadje](#), [Mohammed Ali Al-Garadi](#), [Abeed Sarker](#)

Abstract: **Large...** [▽ More](#)

Submitted 27 March, 2024; originally announced March 2024.

221. [arXiv:2403.16851](#) [pdf] cs.DL cs.AI cs.CL cs.LG doi 10.1002/asi.70028

Can **social media** provide early warning of retraction? Evidence from critical tweets identified by human **annotation** and **large language models**

Authors: [Er-Te Zheng](#), [Hui-Zhen Fu](#), [Mike Thelwall](#), [Zhichao Fang](#)

Abstract: Timely detection of problematic research is essential for safeguarding scientific integrity. To explore whether **social**... [▽ More](#)

Submitted 25 September, 2025; v1 submitted 25 March, 2024; **originally announced** March 2024.

Comments: 27 pages, 5 figures

Journal ref: Journal of the Association for Information Science and Technology, 2025

222. [arXiv:2403.13272](#) [pdf, other] cs.CY cs.CL cs.SI

Community Needs and Assets: A Computational Analysis of Community Conversations

Authors: [Md Towhidul Absar Chowdhury](#), [Naveen Sharma](#), [Ashiqur R. KhudaBukhsh](#)

Abstract: ...government agencies to quantify the strengths and issues of a community, allowing them to allocate their resources better. Such approaches are transitioning towards leveraging **social**... [▽ More](#)

Submitted 19 March, 2024; **originally announced** March 2024.

223. [arXiv:2403.07678](#) [pdf, other] cs.CL cs.CY doi 10.1145/3677525.3678694

MoralBERT: A Fine-Tuned Language Model for Capturing Moral Values in Social Discussions

Authors: [Vjosa Preniqi](#), [Iacopo Ghinassi](#), [Julia Ive](#), [Charalampos Saitis](#), [Kyriaki Kalimeri](#)

Abstract: ...fine-tuned to capture moral sentiment in social discourse. We describe a framework for both aggregated and domain-adversarial training on multiple heterogeneous MFT human-**annotated** datasets sourced from... [▽ More](#)

Submitted 19 July, 2024; v1 submitted 12 March, 2024; **originally announced** March 2024.

Journal ref: ACM 4th International Conference on Information Technology for Social Good (GoodIT 2024)

224. [arXiv:2403.05216](#) [pdf, other] cs.CL cs.SI

SocialPET: Socially Informed Pattern Exploiting Training for **Few-Shot Stance Detection** in **Social Media**

Authors: [Parisa Jamadi Khiabani](#), [Arkaitz Zubiaga](#)

Abstract: **Stance**... [▽ More](#)

Submitted 8 March, 2024; **originally announced** March 2024.

225. [arXiv:2403.02472](#) [pdf, other] cs.CL doi 10.1109/ACCESS.2024.3513220

OffensiveLang: A Community Based Implicit Offensive Language Dataset

Authors: [Amit Das](#), [Mostafa Rahgouy](#), [Dongji Feng](#), [Zheng Zhang](#), [Tathagata Bhattacharya](#), [Nilanjana Raychawdhary](#), [Fatemeh Jamshidi](#), [Vinija Jain](#), [Aman Chadha](#), [Mary Sandage](#), [Lauramarie Pope](#), [Gerry Dozier](#), [Cheryl Seals](#)

Abstract: The widespread presence of hateful languages on **social**... [▽ More](#)

Submitted 14 December, 2024; v1 submitted 4 March, 2024; **originally announced** March 2024.

Journal ref: in IEEE Access, vol. 12, pp. 185661-185672, 2024

226. [arXiv:2403.01270](#) [pdf] cs.CL

A comprehensive cross-language framework for harmful content detection with the aid of **sentiment analysis**

Authors: [Mohammad Dehghani](#)

Abstract: In today's digital world, **social...** [▽ More](#)

Submitted 2 March, 2024; **originally announced** March 2024.

227. [arXiv:2403.00418](#) [[pdf](#), [other](#)] [cs.CL](#)

LLMs for Targeted Sentiment in News Headlines: Exploring the Descriptive-Prescriptive Dilemma

Authors: [Jana Juroš](#), [Laura Majer](#), [Jan Šnajder](#)

Abstract: News headlines often evoke sentiment by intentionally portraying entities in particular ways, making targeted **sentiment...** [▽ More](#)

Submitted 30 October, 2024; **v1** submitted 1 March, 2024; **originally announced** March 2024.

Comments: Presented at 14th Workshop on Computational Approaches to Subjectivity, Sentiment, & **Social Media** Analysis (WASSA) at ACL 2024

228. [arXiv:2402.18169](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

MIKO: Multimodal Intention Knowledge Distillation from **Large Language Models** for **Social-Media** Commonsense Discovery

Authors: [Feihong Lu](#), [Weiqi Wang](#), [Yangyifei Luo](#), [Ziqin Zhu](#), [Qingyun Sun](#), [Baixuan Xu](#), [Haochen Shi](#), [Shiqi Gao](#), [Qian Li](#), [Yangqiu Song](#), [Jianxin Li](#)

Abstract: **Social...** [▽ More](#)

Submitted 29 February, 2024; **v1** submitted 28 February, 2024; **originally announced** February 2024.

Comments: 11 pages, 5 figures

229. [arXiv:2402.17478](#) [[pdf](#), [other](#)] [cs.CL](#)

Can GPT-4 Identify Propaganda? **Annotation** and Detection of Propaganda Spans in News Articles

Authors: [Maram Hasanain](#), [Fatema Ahmed](#), [Firoj Alam](#)

Abstract: The use of propaganda has spiked on mainstream and **social...** [▽ More](#)

Submitted 27 February, 2024; **originally announced** February 2024.

Comments: Accepted as a full paper at LREC-COLING 2024

230. [arXiv:2402.10948](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [doi](#) [10.1145/3589335.3651584](#).

Zero-shot Explainable Mental Health Analysis on **Social Media** by Incorporating Mental Scales

Authors: [Wenyu Li](#), [Yinuo Zhu](#), [Xin Lin](#), [Ming Li](#), [Ziyue Jiang](#), [Ziqian Zeng](#)

Abstract: Traditional discriminative approaches in mental health analysis are known for their strong capacity but lack interpretability and demand large-scale **annotated** data. The generative approaches, such as those based on... [▽ More](#)

Submitted 14 March, 2024; **v1** submitted 9 February, 2024; **originally announced** February 2024.

Comments: 4 pages, 2 figures

Journal ref: The Web Conference (WWW) 2024, Short Paper

231. [arXiv:2402.10573](#) [[pdf](#), [other](#)] [cs.CL](#)

LinkNER: Linking Local Named Entity Recognition Models to **Large Language Models** using Uncertainty

Authors: [Zhen Zhang](#), [Yuhua Zhao](#), [Hang Gao](#), [Mengting Hu](#)

Abstract: Named Entity Recognition (NER) serves as a fundamental task in natural language understanding, bearing direct implications for web **content**... [▽ More](#)

Submitted 22 December, 2024; **v1** submitted 16 February, 2024; **originally announced** February 2024.

Comments: Accepted by WebConf (WWW'2024), **compared** to the previous version, the code has been released

ACM Class: I.2.7

232. [arXiv:2402.01783](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

Hierarchical Multi-Label Classification of Online Vaccine Concerns

Authors: [Chloe Qinyu Zhu](#), [Rickard Stureborg](#), [Bhuwan Dhingra](#)

Abstract: ...inform the healthcare space by helping public health efforts strategically allocate resources or information campaigns. We explore the task of detecting vaccine concerns in **online**... [▽ More](#)

Submitted 1 February, 2024; **originally announced** February 2024.

Comments: Published in AAAI 2024 Health Intelligence workshop

233. [arXiv:2401.04655](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

DepressionEmo: A novel dataset for multilabel classification of depression emotions

Authors: [Abu Bakar Siddiquir Rahman](#), [Hoang-Thang Ta](#), [Lotfollah Najjar](#), [Azad Azadmanesh](#), [Ali Saffet Gönül](#)

Abstract: ...and impact on individuals. In this paper, we introduce a novel dataset named DepressionEmo designed to detect 8 emotions associated with depression by 6037 examples of long **Reddit** user posts. This dataset was created through a majority vote over inputs by... [▽ More](#)

Submitted 9 January, 2024; **originally announced** January 2024.

Comments: 18 pages

234. [arXiv:2311.18063](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#) [cs.SI](#)

TurkishBERTweet: Fast and Reliable **Large Language Model** for Social Media Analysis

Authors: [Ali Najafi](#), [Onur Varol](#)

Abstract: Turkish is one of the most popular languages in the world. Wide us of this language on **social**... [▽ More](#)

Submitted 29 November, 2023; **originally announced** November 2023.

Comments: 21 pages, 4 figures, 8 tables

235. [arXiv:2311.14759](#) [[pdf](#), [other](#)] [q-fin.ST](#) [cs.LG](#) [stat.ML](#)

Deep Learning and NLP in Cryptocurrency Forecasting: Integrating Financial, Blockchain, and **Social Media** Data

Authors: [Vincent Gurgul](#), [Stefan Lessmann](#), [Wolfgang Karl Härdle](#)

Abstract: ...cryptocurrency price forecasting, leveraging Machine Learning (ML) and Natural Language Processing (NLP) techniques, with a focus on Bitcoin and Ethereum. By analysing news and **social**... [▽ More](#)

Submitted 25 October, 2024; **v1** submitted 23 November, 2023; **originally announced** November 2023.

236. [arXiv:2311.14727](#) [pdf, ps, other] cs.CL cs.LG doi 10.1016/j.knosys.2025.114001

Optimal strategies to perform multilingual analysis of social content for a novel dataset in the tourism domain

Authors: [Maxime Masson](#), [Rodrigo Agerri](#), [Christian Sallaberry](#), [Marie-Noelle Bessagnet](#), [Annig Le Parc Lacayrelle](#), [Philippe Roose](#)

Abstract: The rising influence of **social**... [▽ More](#)

Submitted 2 July, 2025; v1 submitted 20 November, 2023; **originally announced** November 2023.

Journal ref: Knowledge-Based Systems, 114001

237. [arXiv:2311.12323](#) [pdf, other] cs.SI cs.CL cs.LG

Modeling Political Orientation of **Social Media** Posts: An Extended Analysis

Authors: [Sadia Kamal](#), [Brenner Little](#), [Jade Gullic](#), [Trevor Harms](#), [Kristin Olofsson](#), [Arunkumar Bagavathi](#)

Abstract: Developing machine learning models to characterize political polarization on online **social**... [▽ More](#)

Submitted 20 November, 2023; **originally announced** November 2023.

238. [arXiv:2311.09602](#) [pdf, other] cs.CL

Language Models (Mostly) Do Not Consider Emotion Triggers When Predicting Emotion

Authors: [Smriti Singh](#), [Cornelia Caragea](#), [Junyi Jessy Li](#)

Abstract: Situations and events evoke emotions in humans, but to what extent do they inform the prediction of emotion detection models? This work investigates how well human-**annotated** emotion triggers correlate with features that models deemed salient in their prediction of emotions. First, we introduce a novel dataset EmoTrigger, consisting of 900... [▽ More](#)

Submitted 25 March, 2024; v1 submitted 16 November, 2023; **originally announced** November 2023.

Comments: NAACL 2024 Camera Ready

239. [arXiv:2311.06446](#) [pdf, ps, other] cs.CL cs.AI

THOS: A **Benchmark** Dataset for Targeted Hate and Offensive Speech

Authors: [Saad Almohaimeed](#), [Saleh Almohaimeed](#), [Ashfaq Ali Shafin](#), [Bogdan Carbunar](#), [Ladislau Bölöni](#)

Abstract: Detecting harmful content on **social**... [▽ More](#)

Submitted 10 November, 2023; **originally announced** November 2023.

240. [arXiv:2311.04284](#) [pdf, other] cs.CL cs.AI

CRAB: Assessing the Strength of Causal Relationships Between Real-world Events

Authors: [Angelika Romanou](#), [Syrielle Montariol](#), [Debjit Paul](#), [Leo Laugier](#), [Karl Aberer](#), [Antoine Bosselut](#)

Abstract: ...they understand the complexity of the underlying network of causal relationships of events in narratives. In this work, we present CRAB, a new Causal Reasoning Assessment **Benchmark** designed to evaluate causal understanding of events in real-world narratives. CRAB contains fine-grained, contextual causality... [▽ More](#)

Submitted 7 November, 2023; **originally announced** November 2023.

241. [arXiv:2311.00671](#) [pdf, other] cs.CL doi 10.1016/j.inffus.2024.102300

Emotion Detection for Misinformation: A Review

Authors: Zhiwei Liu, Tianlin Zhang, Kailai Yang, Paul Thompson, Zeping Yu, Sophia Ananiadou

Abstract: With the advent of **social**... ▽ More

Submitted 1 November, 2023; originally announced November 2023.

Comments: 30 pages, 11 figures

242. [arXiv:2311.00321](#) [pdf, other] cs.CL

HARE: Explainable Hate Speech Detection with Step-by-Step Reasoning

Authors: Yongjin Yang, Joonkee Kim, Yujin Kim, Namgyu Ho, James Thorne, Se-young Yun

Abstract: With the proliferation of **social**... ▽ More

Submitted 22 November, 2023; v1 submitted 1 November, 2023; originally announced November 2023.

Comments: Findings of EMNLP 2023; The first three authors contribute equally

243. [arXiv:2310.14817](#) [pdf, other] cs.LG

Text2Topic: Multi-Label Text Classification System for Efficient Topic Detection in User Generated Content with Zero-Shot Capabilities

Authors: Fengjun Wang, Moran Beladev, Ofri Kleinfeld, Elina Frayerman, Tal Shachar, Eran Fainman, Karen Lastmann Assaraf, Sarai Mizrachi, Benjamin Wang

Abstract: Multi-label **text**... ▽ More

Submitted 23 October, 2023; originally announced October 2023.

244. [arXiv:2310.11467](#) [pdf, other] cs.SE cs.AI cs.LG

Enhancing Binary Code Comment Quality Classification: Integrating Generative AI for Improved Accuracy

Authors: Rohith Arumugam S, Angel Deborah S

Abstract: ...by integrating generated code and comment pairs, to improve model accuracy. The dataset comprises 9048 pairs of code and comments written in the C programming language, each **annotated** as "Useful" or "Not Useful." Additionally, code and comment pairs are generated using a... ▽ More

Submitted 14 October, 2023; originally announced October 2023.

Comments: 11 pages, 2 figures, 2 tables, Has been accepted for the Information Retrieval in Software Engineering track at **Forum** for Information Retrieval **Evaluation** 2023

245. [arXiv:2310.11113](#) [pdf, other] cs.SE doi 10.1145/3697009

Revisiting Sentiment Analysis for Software Engineering in the Era of Large Language Models

Authors: Ting Zhang, Ivana Clairine Irsan, Ferdian Thung, David Lo

Abstract: ...crucial for the effective development and ongoing maintenance of software systems. For software products, analyzing the sentiment of user feedback, e.g., reviews, comments, and **forum** posts can provide valuable insights into user satisfaction and areas for improvement. This can guide the development of future updates and features. However, accurately identify... ▽ More

Submitted 7 September, 2024; v1 submitted 17 October, 2023; originally announced October 2023.

Comments: Accepted by ACM Transactions on Software Engineering and Methodology

246. [arXiv:2310.09238](#) [pdf] cs.CL

BanglaNLP at BLP-2023 Task 2: **Benchmarking** different Transformer Models for **Sentiment Analysis** of Bangla **Social Media** Posts

Authors: [Saumajit Saha](#), [Albert Nanda](#)

Abstract: ...boasts a rich literary tradition, encompassing diverse dialects and language-specific challenges. Despite its linguistic richness and history, Bangla remains categorized as a **low**-... [▽ More](#)

Submitted 17 October, 2023; v1 submitted 13 October, 2023; **originally announced** October 2023.

Comments: 7 pages, 2 figures, workshop

247. [arXiv:2310.07155](#) [[pdf](#), [other](#)] [cs.CL](#)

"A Tale of Two Movements": Identifying and **Comparing** Perspectives in #BlackLivesMatter and #BlueLivesMatter Movements-related Tweets using Weakly Supervised Graph-based Structured Prediction

Authors: [Shamik Roy](#), [Dan Goldwasser](#)

Abstract: **Social**... [▽ More](#)

Submitted 21 October, 2023; v1 submitted 10 October, 2023; **originally announced** October 2023.

Comments: Findings of EMNLP 2023 (Camera Ready)

248. [arXiv:2310.07086](#) [[pdf](#), [ps](#), [other](#)] [cs.AI](#) [cs.SI](#)

Leveraging **Twitter** Data for **Sentiment Analysis** of Transit User Feedback: An NLP Framework

Authors: [Adway Das](#), [Abhishek Kumar Prajapati](#), [Pengxiang Zhang](#), [Mukund Srinath](#), [Andisheh Ranjbari](#)

Abstract: ...often time-consuming, resource intensive, and costly. In this paper, we propose a novel NLP-based framework that harnesses the vast, abundant, and inexpensive data available on **social**... [▽ More](#)

Submitted 10 October, 2025; v1 submitted 10 October, 2023; **originally announced** October 2023.

Comments: Word Count: 5515 words + 3 table (250 words per table) = 6265 words

249. [arXiv:2309.06503](#) [[pdf](#)] [cs.CL](#) [cs.LG](#) [cs.SI](#)

Leveraging **Large Language Models** and Weak Supervision for **Social Media** data **annotation**: an evaluation using COVID-19 self-reported vaccination tweets

Authors: [Ramya Tekumalla](#), [Juan M. Banda](#)

Abstract: The COVID-19 pandemic has presented significant challenges to the healthcare industry and society as a whole. With the rapid development of COVID-19 vaccines, **social**... [▽ More](#)

Submitted 12 September, 2023; **originally announced** September 2023.

250. [arXiv:2309.05668](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Studying the impacts of pre-training using ChatGPT-generated text on downstream tasks

Authors: [Sarthak Anand](#)

Abstract: In recent times, significant advancements have been witnessed in the field of language models, particularly with the emergence of **Large**... [▽ More](#)

Submitted 2 September, 2023; **originally announced** September 2023.

Comments: Master's thesis

251. [arXiv:2306.14030](#) [pdf, other] cs.CL cs.LG

My Boli: Code-mixed Marathi-English Corpora, Pretrained Language Models and Evaluation **Benchmarks**

Authors: [Tanmay Chavan](#), [Omkar Gokhale](#), [Aditya Kane](#), [Shantanu Patankar](#), [Raviraj Joshi](#)

Abstract: The research on code-mixed data is limited due to the unavailability of dedicated code-mixed datasets and pre-trained language models. In this work, we focus on the **low**-... [▽ More](#)

Submitted 20 July, 2023; v1 submitted 24 June, 2023; **originally announced** June 2023.

252. [arXiv:2306.09390](#) [pdf, other] cs.CL cs.AI cs.LG

ChatGPT for Suicide Risk Assessment on **Social Media**: Quantitative Evaluation of Model Performance, Potentials and Limitations

Authors: [Hamideh Ghanadian](#), [Isar Nejadgholi](#), [Hussein Al Osman](#)

Abstract: This paper presents a novel framework for quantitatively evaluating the interactive ChatGPT model in the context of suicidality assessment from **social**... [▽ More](#)

Submitted 15 June, 2023; **originally announced** June 2023.

253. [arXiv:2306.07499](#) [pdf, other] cs.CL cs.AI cs.LG

Improving Opinion-based Question Answering Systems Through Label Error Detection and Overwrite

Authors: [Xiao Yang](#), [Ahmed K. Mohamed](#), [Shashank Jain](#), [Stanislav Peshterliev](#), [Debojeet Chatterjee](#), [Hanwen Zha](#), [Nikita Bhalla](#), [Gagan Aneja](#), [Pranab Mohanty](#)

Abstract: Label error is a ubiquitous problem in **annotated** data. Large amounts of label error substantially degrades the quality of deep learning models. Existing methods to tackle the label error problem largely focus on the classification task, and either rely on... [▽ More](#)

Submitted 12 June, 2023; **originally announced** June 2023.

254. [arXiv:2306.06147](#) [pdf] cs.CL cs.AI doi 10.1145/3580305.3599904

SentiGOLD: A Large Bangla Gold Standard Multi-Domain **Sentiment Analysis** Dataset and its Evaluation

Authors: [Md. Ekramul Islam](#), [Labib Chowdhury](#), [Faisal Ahamed Khan](#), [Shazzad Hossain](#), [Sourave Hossain](#), [Mohammad Mamun Or Rashid](#), [Nabeel Mohammed](#), [Mohammad Ruhul Amin](#)

Abstract: This study introduces SentiGOLD, a Bangla multi-domain **sentiment**... [▽ More](#)

Submitted 9 June, 2023; **originally announced** June 2023.

Comments: Accepted in KDD 2023 Applied Data Science Track; 12 pages, 14 figures

255. [arXiv:2306.04964](#) [pdf, other] cs.CL cs.LG

Leveraging Language Identification to Enhance Code-Mixed **Text Classification**

Authors: [Gauri Takawane](#), [Abhishek Phaltankar](#), [Varad Patwardhan](#), [Aryan Patil](#), [Raviraj Joshi](#), [Mukta S. Takalikar](#)

Abstract: ...same text is referred to as Code Mixed. It is evident that there is a growing degree of adaption of the use of code-mixed data, especially English with a regional language, on **social**...

[▽ More](#)

Submitted 8 June, 2023; **originally announced** June 2023.

256. [arXiv:2305.15722](#) [pdf, other] cs.CL cs.LG doi 10.1109/I2CT57861.2023.10126273

Comparative Study of Pre-Trained BERT Models for Code-Mixed Hindi-English Data

Authors: [Aryan Patil](#), [Varad Patwardhan](#), [Abhishek Phaltankar](#), [Gauri Takawane](#), [Raviraj Joshi](#)

Abstract: The term "Code Mixed" refers to the use of more than one language in the same text. This phenomenon is predominantly observed on [social...](#) [▽ More](#)

Submitted 26 May, 2023; **v1** submitted 25 May, 2023; **originally announced** May 2023.

Comments: Accepted at IEEE 8th International Conference for Convergence in Technology

257. [arXiv:2305.14291](#) [[pdf](#), [other](#)] [cs.CL](#)

Evaluation of African American Language Bias in Natural Language Generation

Authors: [Nicholas Deas](#), [Jessi Grieser](#), [Shana Kleiner](#), [Desmond Patton](#), [Elsbeth Turcan](#), [Kathleen McKeown](#)

Abstract: We evaluate how well LLMs understand African American Language (AAL) in [comparison](#) to their performance on White Mainstream English (WME), the encouraged "standard" form of English taught in American classrooms. We measure LLM performance using automatic metrics and human judgments for two tasks: a counterpart generation task, where a model generates... [▽ More](#)

Submitted 12 November, 2023; **v1** submitted 23 May, 2023; **originally announced** May 2023.

Comments: EMNLP 2023 Camera-Ready

258. [arXiv:2305.12692](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

MetaAdapt: Domain Adaptive [Few-Shot](#) Misinformation Detection via Meta Learning

Authors: [Zhenrui Yue](#), [Huimin Zeng](#), [Yang Zhang](#), [Lanyu Shang](#), [Dong Wang](#)

Abstract: With emerging topics (e.g., COVID-19) on [social...](#) [▽ More](#)

Submitted 22 May, 2023; **originally announced** May 2023.

Comments: Accepted to ACL 2023

259. [arXiv:2305.10971](#) [[pdf](#), [other](#)] [cs.CL](#)

NollySenti: Leveraging Transfer Learning and Machine Translation for Nigerian Movie [Sentiment Classification](#)

Authors: [Iyanuoluwa Shode](#), [David Ifeoluwa Adelani](#), [Jing Peng](#), [Anna Feldman](#)

Abstract: ...labeled corpora for African languages. However, they are often available in a single domain and may not generalize to other domains. In this paper, we focus on the task of [sentiment...](#) [▽ More](#)

Submitted 22 August, 2023; **v1** submitted 18 May, 2023; **originally announced** May 2023.

Comments: Accepted to ACL 2023 (main conference)

260. [arXiv:2304.13634](#) [[pdf](#), [other](#)] [cs.CL](#)

HausaNLP at SemEval-2023 Task 12: Leveraging African [Low Resource](#) TweetData for [Sentiment Analysis](#)

Authors: [Saheed Abdullahi Salahudeen](#), [Falalu Ibrahim Lawan](#), [Ahmad Mustapha Wali](#), [Amina Abubakar Imam](#), [Aliyu Rabi'u Shuaibu](#), [Aliyu Yusuf](#), [Nur Bala Rabi'u](#), [Musa Bello](#), [Shamsuddeen Umaru Adamu](#), [Saminu Mohammad Aliyu](#), [Murja Sani Gadanya](#), [Sanah Abdullahi Muaz](#), [Mahmoud Said Ahmad](#), [Abdulkadir Abdullahi](#), [Abdulmalik Yusuf Jamoh](#)

Abstract: We present the findings of SemEval-2023 Task 12, a shared task on [sentiment...](#) [▽ More](#)

Submitted 26 April, 2023; originally announced April 2023.

261. [arXiv:2304.06588](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.SI](#)

ChatGPT-4 Outperforms Experts and Crowd Workers in **Annotating** Political **Twitter** Messages with **Zero-Shot** Learning

Authors: [Petter Törnberg](#)

Abstract: This paper assesses the accuracy, reliability and bias of the **Large...** [▽ More](#)

Submitted 13 April, 2023; originally announced April 2023.

Comments: 5 pages, 3 figures

262. [arXiv:2303.15016](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.IR](#) [cs.MM](#)

Borrowing Human Senses: Comment-Aware Self-Training for **Social Media** Multimodal Classification

Authors: [Chunpu Xu](#), [Jing Li](#)

Abstract: **Social...** [▽ More](#)

Submitted 27 March, 2023; originally announced March 2023.

Comments: accepted to EMNLP 2022

263. [arXiv:2303.08032](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#) [doi](#) [10.1017/nlp.2024.54](#)

Verifying the Robustness of Automatic Credibility Assessment

Authors: [Piotr Przybyła](#), [Alexander Shvets](#), [Horacio Saggion](#)

Abstract: **Text...** [▽ More](#)

Submitted 21 November, 2024; v1 submitted 14 March, 2023; originally announced March 2023.

Journal ref: Nat. lang. process. 31 (2025) 1134-1162

264. [arXiv:2302.13619](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Orca: A **Few-shot Benchmark** for Chinese Conversational Machine Reading Comprehension

Authors: [Nuo Chen](#), [Hongguang Li](#), [Junqing He](#), [Yinan Bao](#), [Xinshi Lin](#), [Qi Yang](#), [Jianfeng Liu](#), [Ruyi Gan](#), [Jiaxing Zhang](#), [Baoyuan Wang](#), [Jia Li](#)

Abstract: ...(CMRC) task aims to answer questions in conversations, which has been a hot research topic in recent years because of its wide applications. However, existing CMRC **benchmarks** in which each conversation is assigned a static passage are inconsistent with real scenarios. Thus, model's comprehension ability towards real scenarios are hard to evaluate reasona... [▽ More](#)

Submitted 13 October, 2023; v1 submitted 27 February, 2023; originally announced February 2023.

Comments: 14 pages

Journal ref: EMNLP 2023

265. [arXiv:2301.08774](#) [[pdf](#), [other](#)] [cs.SI](#) [cs.AI](#)

DoubleH: **Twitter** User **Stance Detection** via Bipartite Graph Neural Networks

Authors: [Chong Zhang](#), [Zhenkun Zhou](#), [Xingyu Peng](#), [Ke Xu](#)

Abstract: Given the development and abundance of **social...** [▽ More](#)

Submitted 20 January, 2023; originally announced January 2023.

266. [arXiv:2301.04535](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.SI](#)

Few-shot Learning for Cross-Target Stance Detection by Aggregating Multimodal Embeddings

Authors: [Parisa Jamadi Khiabani](#), [Arkaitz Zubiaga](#)

Abstract: Despite the increasing popularity of the **stance**... [▽ More](#)

Submitted 31 March, 2023; v1 submitted 11 January, 2023; **originally announced** January 2023.

Comments: To appear in IEEE Transactions on Computational Social Systems

267. [arXiv:2212.10039](#) [[pdf](#), [other](#)] [cs.CL](#)

A Twitter BERT Approach for Offensive Language Detection in Marathi

Authors: [Tanmay Chavan](#), [Shantanu Patankar](#), [Aditya Kane](#), [Omkar Gokhale](#), [Raviraj Joshi](#)

Abstract: Automated offensive language detection is essential in combating the spread of hate speech, particularly in **social**... [▽ More](#)

Submitted 20 December, 2022; **originally announced** December 2022.

268. [arXiv:2212.00851](#) [[pdf](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#) [cs.SI](#) [doi](#) [10.1007/s10579-024-09723-1](#)

SOLD: Sinhala Offensive Language Dataset

Authors: [Tharindu Ranasinghe](#), [Isuri Anuradha](#), [Damith Premasiri](#), [Kanishka Silva](#), [Hansi Hettiarachchi](#), [Lasitha Uyangodage](#), [Marcos Zampieri](#)

Abstract: ...natural language processing (NLP) communities, motivating the development of various systems trained to detect potentially harmful content automatically. These systems require **annotated** datasets to train the machine learning (ML) models. However, with a few notable exceptions, most datasets on this topic have dealt with English and a few other high-resource... [▽ More](#)

Submitted 28 March, 2024; v1 submitted 1 December, 2022; **originally announced** December 2022.

Comments: Accepted to Language Resources and **Evaluation**, Springer

269. [arXiv:2212.00668](#) [[pdf](#), [other](#)] [cs.CV](#)

Leveraging Large-scale Multimedia Datasets to Refine Content Moderation Models

Authors: [Ioannis Sarridis](#), [Christos Koutlis](#), [Olga Papadopoulou](#), [Symeon Papadopoulos](#)

Abstract: The sheer volume of online **user**-... [▽ More](#)

Submitted 1 December, 2022; **originally announced** December 2022.

270. [arXiv:2211.11798](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Can You Label Less by Using Out-of-Domain Data? Active & Transfer Learning with Few-shot Instructions

Authors: [Rafal Kocielnik](#), [Sara Kangaslahti](#), [Shrimai Prabhumoye](#), [Meena Hari](#), [R. Michael Alvarez](#), [Anima Anandkumar](#)

Abstract: Labeling **social**-... [▽ More](#)

Submitted 21 November, 2022; **originally announced** November 2022.

Comments: Accepted to NeurIPS Workshop on Transfer Learning for Natural Language Processing, 2022, New Orleans

271. [arXiv:2210.15616](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.IR](#) [cs.LG](#)

Cross-Domain Neural Entity Linking

Authors: [Hassan Soliman](#)

Abstract: Entity Linking is the task of matching a mention to an entity in a given knowledge base (KB). It contributes to **annotating** a massive amount of documents existing on the Web to harness new facts about their matched entities. However, existing Entity Linking systems focus on developing models that are typically domain-dependent and robust only to a particular... [▽ More](#)

Submitted 28 September, 2022; **originally announced** October 2022.

Comments: Master's thesis, 76 pages, 26 figures

272. [arXiv:2210.05715](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Language Independent **Stance Detection**: Social Interaction-based Embeddings and **Large Language Models**

Authors: [Joseba Fernandez de Landa](#), [Rodrigo Agerri](#)

Abstract: The large majority of the research performed on **stance**... [▽ More](#)

Submitted 27 February, 2025; **v1** submitted 11 October, 2022; **originally announced** October 2022.

273. [arXiv:2210.05425](#) [[pdf](#)] [cs.CL](#) [cs.AI](#) [cs.CY](#) [cs.LG](#)

COVID-19-related Nepali Tweets Classification in a **Low Resource** Setting

Authors: [Rabin Adhikari](#), [Safal Thapaliya](#), [Nirajan Basnet](#), [Samip Poudel](#), [Aman Shakya](#), [Bishesh Khanal](#)

Abstract: Billions of people across the globe have been using **social**... [▽ More](#)

Submitted 11 October, 2022; **originally announced** October 2022.

Comments: Accepted at the 7th **Social Media** Mining for Health (#SMM4H) Workshop, co-located at Coling 2022

274. [arXiv:2210.03235](#) [[pdf](#), [other](#)] [cs.CL](#)

Improving Large-scale Paraphrase Acquisition and Generation

Authors: [Yao Dou](#), [Chao Jiang](#), [Wei Xu](#)

Abstract: This paper addresses the quality issues in existing **Twitter**-based paraphrase datasets, and discusses the necessity of using two separate definitions of paraphrase for identification and generation tasks. We present a new Multi-Topic Paraphrase in... [▽ More](#)

Submitted 7 November, 2022; **v1** submitted 6 October, 2022; **originally announced** October 2022.

Comments: The project webpage is at <http://twitter-paraphrase.com/> Accepted at EMNLP 2022

275. [arXiv:2209.10482](#) [[pdf](#), [other](#)] [cs.CL](#)

SMTCE: A **Social Media Text Classification** Evaluation **Benchmark** and BERTology Models for Vietnamese

Authors: [Luan Thanh Nguyen](#), [Kiet Van Nguyen](#), [Ngan Luu-Thuy Nguyen](#)

Abstract: **Text**... [▽ More](#)

Submitted 21 September, 2022; **originally announced** September 2022.

Comments: Accepted at The 36th annual Meeting of Pacific Asia Conference on Language, Information and Computation (PACLIC 36)

276. [arXiv:2209.08207](#) [[pdf](#), [other](#)] [cs.CL](#)

APPDIA: A Discourse-aware Transformer-based Style Transfer Model for Offensive **Social Media** Conversations

Authors: [Katherine Atwell](#), [Sabit Hassan](#), [Malihe Alikhani](#)

Abstract: Using style-transfer models to reduce offensiveness of **social**... [▽ More](#)

Submitted 16 September, 2022; **originally announced** September 2022.

277. [arXiv:2209.02765](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

Depression Symptoms Modelling from **Social Media** Text: A Semi-supervised Learning Approach

Authors: [Nawshad Farruque](#), [Randy Goebel](#), [Sudhakar Sivapalan](#), [Osmar Zaiane](#)

Abstract: A fundamental component of user-level **social**... [▽ More](#)

Submitted 28 September, 2022; **v1** submitted 6 September, 2022; **originally announced** September 2022.

Comments: Title and relevant changes are made

278. [arXiv:2208.09021](#) [[pdf](#), [other](#)] [cs.CV](#) [cs.AI](#) [cs.CL](#) [cs.LG](#)

VAuLT: Augmenting the Vision-and-Language Transformer for **Sentiment Classification** on **Social Media**

Authors: [Georgios Chochlakis](#), [Tejas Srinivasan](#), [Jesse Thomason](#), [Shrikanth Narayanan](#)

Abstract: ...where the language input is simple, literal, and descriptive, therefore lacking linguistic diversity. So, when working with multimedia data in the wild, such as multimodal **social**... [▽ More](#)

Submitted 25 January, 2023; **v1** submitted 18 August, 2022; **originally announced** August 2022.

Comments: 5 pages, 1 figure

279. [arXiv:2208.05545](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CY](#) [cs.LG](#)

The Moral Foundations **Reddit** Corpus

Authors: [Jackson Trager](#), [Alireza S. Ziabari](#), [Elnaz Rahmati](#), [Aida Mostafazadeh Davani](#), [Prezi Golazizian](#), [Farzan Karimi-Malekabadi](#), [Ali Omrani](#), [Zhihe Li](#), [Brendan Kennedy](#), [Georgios Chochlakis](#), [Nils Karl Reimer](#), [Melissa Reyes](#), [Kelsey Cheng](#), [Mellow Wei](#), [Christina Merrifield](#), [Arta Khosravi](#), [Evans Alvarez](#), [Morteza Dehghani](#)

Abstract: ...in Natural Language Processing (NLP) have been used to detect moral sentiment from textual data, but achieving strong performance in such subjective tasks requires large, hand-**annotated** datasets. Previous corpora... [▽ More](#)

Submitted 17 March, 2026; **v1** submitted 10 August, 2022; **originally announced** August 2022.

280. [arXiv:2207.13005](#) [[pdf](#), [other](#)] [cs.CL](#)

Hansel: A Chinese **Few-Shot** and **Zero-Shot** Entity Linking **Benchmark**

Authors: [Zhenran Xu](#), [Zifei Shan](#), [Yuxin Li](#), [Baotian Hu](#), [Bing Qin](#)

Abstract: ...Linking (EL) systems entrench a popularity bias, yet there is no dataset focusing on tail and emerging entities in languages other than English. We present Hansel, a new **benchmark** in Chinese that fills the vacancy of non-English... [▽ More](#)

Submitted 29 October, 2023; **v1** submitted 26 July, 2022; **originally announced** July 2022.

Comments: WSDM 2023

281. [arXiv:2204.10710](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.AI](#)

Tweets2Stance: Users **stance detection** exploiting **Zero-Shot** Learning Algorithms on Tweets

Authors: [Margherita Gambini](#), [Tiziano Fagni](#), [Caterina Senette](#), [Maurizio Tesconi](#)

Abstract: In the last years there has been a growing attention towards predicting the political orientation of active **social**... [▽ More](#)

Submitted 22 April, 2022; originally announced April 2022.

282. [arXiv:2204.08398](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

L3Cube-HingCorpus and HingBERT: A Code Mixed Hindi-English Dataset and BERT Language Models

Authors: [Ravindra Nayak](#), [Raviraj Joshi](#)

Abstract: Code-switching occurs when more than one language is mixed in a given sentence or a conversation. This phenomenon is more prominent on [social...](#) [▽ More](#)

Submitted 18 April, 2022; originally announced April 2022.

283. [arXiv:2204.07010](#) [[pdf](#), [other](#)] [cs.CL](#)

Anti-Asian Hate Speech Detection via Data Augmented Semantic Relation Inference

Authors: [Jiaxuan Li](#), [Yue Ning](#)

Abstract: With the spreading of hate speech on [social...](#) [▽ More](#)

Submitted 14 April, 2022; originally announced April 2022.

Comments: To appear in Proceedings of the 16th International AAAI Conference on Web and [Social Media](#) (ICWSM)

284. [arXiv:2204.03839](#) [[pdf](#), [other](#)] [cs.CL](#)

Infusing Knowledge from Wikipedia to Enhance [Stance Detection](#)

Authors: [Zihao He](#), [Negar Mokherberian](#), [Kristina Lerman](#)

Abstract: [Stance...](#) [▽ More](#)

Submitted 8 April, 2022; originally announced April 2022.

285. [arXiv:2203.15936](#) [[pdf](#), [other](#)] [cs.LG](#)

Supervised Graph Contrastive Learning for [Few-shot](#) Node Classification

Authors: [Zhen Tan](#), [Kaize Ding](#), [Ruocheng Guo](#), [Huan Liu](#)

Abstract: ...are present in many real-world applications, such as financial fraud detection, commercial recommendation, and social network analysis. But given the high cost of graph [annotation](#) or labeling, we face a severe graph label-scarcity problem, i.e., a graph might have a few labeled nodes. One example of such a problem is the so-called \textit{...} [▽ More](#)

Submitted 5 August, 2022; v1 submitted 29 March, 2022; originally announced March 2022.

Comments: Accepted to ECML-PKDD 2022

286. [arXiv:2203.14267](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

bitsa_nlp@LT-EDI-ACL2022: Leveraging Pretrained Language Models for Detecting Homophobia and Transphobia in [Social Media](#) Comments

Authors: [Vitthal Bhandari](#), [Poonam Goyal](#)

Abstract: Online social networks are ubiquitous and user-friendly. Nevertheless, it is vital to detect and moderate offensive content to maintain decency and empathy. However, mining [social...](#)

[▽ More](#)

Submitted 9 April, 2022; v1 submitted 27 March, 2022; originally announced March 2022.

Comments: 6 pages, Accepted at LT-EDI workshop ACL 2022. Camera ready version. Addressed all reviewer comments. Added Baseline methods and Ablation study

287. [arXiv:2110.10318](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#)

Improved Multilingual Language Model Pretraining for **Social Media** Text via Translation Pair Prediction

Authors: [Shubhanshu Mishra](#), [Aria Haghighi](#)

Abstract: We evaluate a simple approach to improving **zero**-... [▽ More](#)

Submitted 19 October, 2021; **originally announced** October 2021.

Comments: Camera ready version. Accepted to WNUT 2021. Code for reproducing the experiments can be found at: <https://github.com/twitter-research/multilingual-alignment-tpp>

MSC Class: 68T50; 68T07 **ACM Class:** I.2.7

288. [arXiv:2110.05021](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#)

Cross Domain Emotion Recognition using **Few Shot** Knowledge Transfer

Authors: [Justin Olah](#), [Sabyasachee Baruah](#), [Digbalay Bose](#), [Shrikanth Narayanan](#)

Abstract: Emotion recognition from text is a challenging task due to diverse emotion taxonomies, lack of reliable labeled data in different domains, and highly subjective **annotation** standards....

[▽ More](#)

Submitted 11 October, 2021; **originally announced** October 2021.

Comments: 5 pages, 4 figures

289. [arXiv:2109.12773](#) [[pdf](#), [other](#)] [cs.CL](#)

Rumour Detection via **Zero-shot** Cross-lingual Transfer Learning

Authors: [Lin Tian](#), [Xiuzhen Zhang](#), [Jey Han Lau](#)

Abstract: Most rumour detection models for **social**... [▽ More](#)

Submitted 26 September, 2021; **originally announced** September 2021.

Comments: ECML-PKDD 2021

290. [arXiv:2109.06050](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#)

Few-Shot Cross-Lingual **Stance Detection** with Sentiment-Based Pre-Training

Authors: [Momchil Hardalov](#), [Arnav Arora](#), [Preslav Nakov](#), [Isabelle Augenstein](#)

Abstract: The goal of **stance**... [▽ More](#)

Submitted 21 December, 2021; **v1** submitted 13 September, 2021; **originally announced** September 2021.

Comments: Accepted to AAAI 2022 (Preprint version)

291. [arXiv:2108.03805](#) [[pdf](#), [other](#)] [cs.SI](#)

Learning to Detect **Few-Shot-Few-Clue** Misinformation

Authors: [Qiang Zhang](#), [Hongbin Huang](#), [Shangsong Liang](#), [Zaiqiao Meng](#), [Emine Yilmaz](#)

Abstract: ...disquieting due to the lack of careful manual review. Consequently, a large volume of false textual information has been disseminating for a long time since the prevalence of **social**...

[▽ More](#)

Submitted 9 August, 2021; **originally announced** August 2021.

292. [arXiv:2105.05996](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#) [cs.SI](#)

Multilingual Offensive Language Identification for **Low-resource** Languages

Authors: [Tharindu Ranasinghe](#), [Marcos Zampieri](#)

Abstract: Offensive content is pervasive in **social**... [▽ More](#)

Submitted 20 May, 2021; **v1** submitted 12 May, 2021; **originally announced** May 2021.

Comments: Accepted to ACM Transactions on Asian and **Low-Resource** Language Information Processing (TALLIP). This is an extended version of a paper accepted to EMNLP. arXiv admin note: substantial text overlap with arXiv:2010.05324

293. [arXiv:2104.12250](#) [[pdf](#), [other](#)] cs.CL

XLM-T: Multilingual Language Models in **Twitter** for **Sentiment Analysis** and Beyond

Authors: [Francesco Barbieri](#), [Luis Espinosa Anke](#), [Jose Camacho-Collados](#)

Abstract: ...their multilingual capacity has recently attracted considerable attention. However, current analyses have almost exclusively focused on (multilingual variants of) standard **benchmarks**, and have relied on clean pre-training and... [▽ More](#)

Submitted 11 May, 2022; **v1** submitted 25 April, 2021; **originally announced** April 2021.

Comments: LREC 2022. Code and data available at <https://github.com/cardiffnlp/xlm-t>

294. [arXiv:2103.14797](#) [[pdf](#), [other](#)] cs.CL cs.LG

Unsupervised Self-Training for **Sentiment Analysis** of Code-Switched Data

Authors: [Akshat Gupta](#), [Sargam Menghani](#), [Sai Krishna Rallabandi](#), [Alan W Black](#)

Abstract: **Sentiment...** [▽ More](#)

Submitted 1 October, 2021; **v1** submitted 26 March, 2021; **originally announced** March 2021.

295. [arXiv:2103.05149](#) [[pdf](#), [ps](#), [other](#)] cs.CL

Contrastive Semi-supervised Learning for ASR

Authors: [Alex Xiao](#), [Christian Fuegen](#), [Abdelrahman Mohamed](#)

Abstract: ...the most adopted method for pre-training automatic speech recognition (ASR) models. However, its performance suffers from the supervised teacher model's degrading quality in **low...** [▽ More](#)

Submitted 8 March, 2021; **originally announced** March 2021.

296. [arXiv:2010.06906](#) [[pdf](#), [other](#)] cs.CL cs.LG cs.SI

No Rumours Please! A Multi-Indic-Lingual Approach for COVID Fake-Tweet Detection

Authors: [Debanjana Kar](#), [Mohit Bhardwaj](#), [Suranjana Samanta](#), [Amar Prakash Azad](#)

Abstract: ...widespread menace created by the present global pandemic COVID-19 has had an unprecedented effect on our lives. Man-kind is going through humongous fear and dependence on **social...** [▽ More](#)

Submitted 14 October, 2020; **originally announced** October 2020.

Comments: 6 pages, 4 figures

297. [arXiv:2007.12544](#) [[pdf](#), [other](#)] cs.CL

FiSSA at SemEval-2020 Task 9: Fine-tuned For Feelings

Authors: [Bertelt Braaksma](#), [Richard Scholtens](#), [Stan van Suijlekom](#), [Remy Wang](#), [Ahmet Üstün](#)

Abstract: In this paper, we present our approach for **sentiment...** [▽ More](#)

Submitted 19 October, 2020; **v1** submitted 24 July, 2020; **originally announced** July 2020.

Comments: In Proceedings of the 14th International Workshop on Semantic **Evaluation** (SemEval-2020), Barcelona, Spain, December. Association for Computational Linguistics

298. [arXiv:2006.00206](#) [[pdf](#), [other](#)] cs.CL

Corpus Creation for **Sentiment Analysis** in Code-Mixed Tamil-English Text

Authors: [Bharathi Raja Chakravarthi](#), [Vigneshwaran Muralidaran](#), [Ruba Priyadharshini](#), [John P. McCrae](#)

Abstract: Understanding the sentiment of a comment from a video or an image is an essential task in many applications. **Sentiment...** [▽ More](#)

Submitted 30 May, 2020; **originally announced** June 2020.

Journal ref: Proceedings of the 1st Joint Workshop on Spoken Language Technologies for Under-resourced languages (SLTU) and Collaboration and Computing for Under-Resourced Languages (CCURL) 2020

299. [arXiv:2005.01151](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.LG](#)

Let Me Choose: From **Verbal** Context to Font Selection

Authors: [Amirreza Shirani](#), [Franck Dernoncourt](#), [Jose Echevarria](#), [Paul Asente](#), [Nedim Lipka](#), [Thamar Solorio](#)

Abstract: In this paper, we aim to learn associations between visual attributes of fonts and the **verbal** context of the texts they are typically applied to.... [▽ More](#)

Submitted 3 May, 2020; **originally announced** May 2020.

Comments: Accepted to ACL 2020

300. [arXiv:1904.09678](#) [[pdf](#), [other](#)] [cs.CL](#)

UniSent: Universal Adaptable Sentiment Lexica for 1000+ Languages

Authors: [Ehsaneddin Asgari](#), [Fabienne Braune](#), [Benjamin Roth](#), [Christoph Ringlstetter](#), [Mohammad R. K. Mofrad](#)

Abstract: ...languages. Sentiment lexica are vital for **sentiment...** [▽ More](#)

Submitted 28 November, 2019; **v1** submitted 21 April, 2019; **originally announced** April 2019.

301. [arXiv:1903.09878](#) [[pdf](#), [other](#)] [cs.CL](#)

Expanding the **Text Classification** Toolbox with Cross-Lingual Embeddings

Authors: [Meryem M'hamdi](#), [Robert West](#), [Andreea Hossmann](#), [Michael Baeriswyl](#), [Claudiu Musat](#)

Abstract: Most work in **text...** [▽ More](#)

Submitted 26 March, 2019; **v1** submitted 23 March, 2019; **originally announced** March 2019.

[Previous](#)

[1](#)

[2](#)

We gratefully acknowledge support from our **major funders**, **member institutions**, and all contributors.

[About](#) · [Help](#) · [Contact](#) · [Subscribe](#) · [Copyright](#) · [Privacy](#) · [Accessibility](#) · [Operational Status](#)

Major funding support from

SIM NS
FOUNDATION



**SIMONS
FOUNDATION
INTERNATIONAL**


Schmidt Sciences