




arXiv is now an independent nonprofit! [Learn more](#)

arxiv





[Search](#) [Submit](#) [Donate](#) [Log in](#)

Showing 1–200 of 301 results

Search v0.5.6 released 2020-02-24

Query: order: -announced_date_first; size: 200; date_range: from 2019-01-01 to 2026-09-20; include_cross_list: True; terms: AND all="zero-shot" OR "few-shot" OR "low-resource" OR "limited labelled data" OR "limited labeled data" OR "large language model*" OR "prompt-based" OR verbalizer OR "fine-tun*" OR "task-specific" OR "natural language inference"; AND all="text classification" OR "stance detection" OR "sentiment classification" OR "sentiment analysis" OR "content analysis" OR annotation; AND all="social media" OR reddit OR twitter OR "user-generated content" OR

[Simple Search](#)

"online discourse" OR forum*; AND all=compar* OR benchmark*
OR evaluat*

Refine query

New search

200



results per page. Sort results by

Announcement date (newest first)



Go

1

2

Next

1. [arXiv:2609.15205](#) [[pdf](#), [ps](#), [other](#)] [cs.DB](#) [cs.AI](#)

TEAR: Table Extraction with Attribute Recommendation from Texts via **Large Language Models**

Authors: [Tong Li](#), [Shuye Ding](#), [Jiachuan Wang](#), [Yongqi Zhang](#), [Shuangyin Li](#), [Lei Chen](#), [Bo Li](#)

Abstract: Table extraction from texts is an important task for information systems, and recent approaches that prompt **large...** [▽ More](#)

Submitted 14 September, 2026; **originally announced** September 2026.

Comments: Preprint version of the paper accepted to SIGMOD 2027 Round 2

2. [arXiv:2609.12653](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

SWARM: A Multilingual Human-**Annotated** Dataset for Russian Propaganda Detection in Search Engine Results

Authors: [Manuel Tonneau](#), [Abhinav Dubey](#), [Farhan Shaikh](#), [Ilaria Vitulano](#), [Martha Stolze](#), [Hale Dedeoglu](#), [Clara Riechert](#), [Ella Kuka](#), [Maryna Sydorova](#), [Mykola Makhortykh](#), [Elizaveta Kuznetsova](#)

Abstract: Russian state propaganda spreads across many languages and online spaces. Yet, most computational work examines only one such space, usually **social...** [▽ More](#)

Submitted 11 September, 2026; **originally announced** September 2026.

Comments: Accepted at the 10th Workshop on Online Abuse and Harms (WOAH 2026)

3. [arXiv:2609.08156](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

When Metrics Reward the Worst Translations: Internalizing Cultural Reasoning for **Social Media** Translation Evaluation

Authors: [Yiwen Qiu](#), [Linjuan Wu](#), [Dingming Li](#), [Yizhou Liu](#), [Zixuan Wang](#), [Haolei Xu](#), [Ye Guo](#), [Daoxin Zhang](#), [Weiming Lu](#), [Yongliang Shen](#)

Abstract: Automatic translation quality metrics trained on general-domain corpora systematically fail on **social...** [▽ More](#)

Submitted 7 September, 2026; **originally announced** September 2026.

4. [arXiv:2609.07766](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Bag of Tricks or Bag of Myths? Reducing Modeling Complexity with Task Knowledge in Explainable Suicide Risk Assessment

Authors: [Shlok Shelat](#), [Shrey Salvi](#), [Souvik Roy](#), [Manas Gaur](#), [Amit Sheth](#)

Abstract: Assessing suicide risk from **social...** [▽ More](#)

Submitted 7 September, 2026; **originally announced** September 2026.

Comments: This work was shaped by the IEEE Big Data Cup, where our team, NIDHI, ranked 3rd overall. NIDHI (Neurosymbolic Intelligent Diagnostics for Holistic Trauma-Informed Care) is an effort from IAIRO.ai . 21 pages, 6 Tables, 4 Figures

5. [arXiv:2609.06963](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

TurEngMix: A Text Corpus and **Benchmark** for Turkish-English Code-Mixed Language Identification and Named Entity Recognition

Authors: [Ilayda Dogan](#), [Phuong-Anh Nguyen-Le](#), [Julia Mendelsohn](#)

Abstract: Natural language processing systems underperform on code-mixed text, particularly for **low**... [▽ More](#)

Submitted 6 September, 2026; **originally announced** September 2026.

Comments: Accepted to W-NUT 2026 (11th Workshop on Natural User-generated Text), co-located with EMNLP 2026. 15 pages, 4 figures, 12 tables

6. [arXiv:2609.00206](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.AI](#)

Distributed Implicit Harm: A Compositional Safety Blind Spot in MLLM-Based Video Moderation

Authors: [Ruotong Wang](#), [Zihao Zhu](#), [Siwei Lyu](#), [Xin Tao](#), [Baoyuan Wu](#)

Abstract: Despite their growing use in video moderation, multimodal **large**... [▽ More](#)

Submitted 31 August, 2026; **originally announced** September 2026.

7. [arXiv:2608.30903](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

MMDS-Bench: **Benchmarking** Multimodal **Large Language Models** on Dynamic Stance in **Social Media** Interactions

Authors: [Yuzhe Ding](#), [Kang He](#), [Li Zheng](#), [Shengwu Zheng](#), [Teng Shi](#), [Fei Li](#), [Chong Teng](#), [Donghong Ji](#)

Abstract: ...how a reply responds to its direct parent message, rather than how a post relates to a fixed topic. Existing work has mainly studied this problem in text-only settings, while **social**... [▽ More](#)

Submitted 31 August, 2026; **originally announced** August 2026.

Comments: Accepted by EMNLP 2026

8. [arXiv:2608.23421](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

A Comprehensive Analysis of Arabic Natural Language Processing Research: Trends, Topic Evolution, and Research Gaps -- A Bibliometric and Topic-Based Study

Authors: [Mullosharaf K. Arabov](#)

Abstract: Arabic Natural Language Processing (NLP) has grown rapidly over the past decade, driven by digital transformation in the Arab world, **social**... [▽ More](#)

Submitted 25 August, 2026; **v1 submitted** 24 August, 2026; **originally announced** August 2026.

Comments: 38 pages, 16 tables, 10 figures. Preprint

9. [arXiv:2608.21385](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CY](#) [cs.SI](#)

A **Social Media** Analysis of Discourse on the Israel--Palestine Conflict on Telegram

Authors: [Michail Zafeiropoulos](#), [Despoina Antonakaki](#), [Sotiris Ioannidis](#)

Abstract: **Social**... [▽ More](#)

Submitted 17 July, 2026; **originally announced** August 2026.

10. [arXiv:2608.20983](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.HC](#) [doi](#) [10.1145/3834580.3838746](#)

Beyond Truth Discovery: A Two-Stage Framework to Assess the Severity of False Claim during Disasters

Authors: [Ruichen Yao](#), [Tejna Dasari](#), [Gulshat Baispay](#), [Aizhan Zaurbek](#), [Yifan Liu](#), [Yaokun Liu](#), [Zelin Li](#), [Dong Wang](#)

Abstract: False information spreads rapidly on **social...** [▽ More](#)

Submitted 30 August, 2026; v1 submitted 21 August, 2026; **originally announced** August 2026.

Comments: The 2026 ACM Conference on Human-AI Complementarity and Alignment (HCOMP 26)

11. [arXiv:2608.18142](#) [[pdf](#)] [cs.AI](#) [cs.CL](#)

Efficient Adaptation of LLMs for Hate Speech Detection in **Low-Resource** Languages: A **Comparative** Study on Roman Urdu

Authors: [Toneema Zubair](#), [Muhammad Junaid Asif](#), [Faisal Kamiran](#), [Hafiz Hassan Saeed](#), [Rana Fayyaz Ahmad](#)

Abstract: It is challenging to detect hate speech in **Low...** [▽ More](#)

Submitted 10 September, 2026; v1 submitted 6 August, 2026; **originally announced** August 2026.

12. [arXiv:2608.14418](#) [[pdf](#), [ps](#), [other](#)] [cs.CR](#)

STINER: Automated Extraction of Strategic Cyber Threat Intelligence from X

Authors: [Yasir Ech-Chammakhy](#), [Oussama Azrara](#), [Jaafar Chbili](#), [Anas Motii](#)

Abstract: ...on high-level insights, such as identifying targeted industries, attributing attacks to specific ransomware groups, and assessing the scale of data loss. Today, X (formerly **Twitter**) has become the fastest source for this intelligence, often hosting real-time breach announcements days before formal vendor reports. Converting this raw chatter into actionable i... [▽ More](#)

Submitted 14 August, 2026; **originally announced** August 2026.

Comments: Accepted at RAID 2026 (29th International Symposium on Research in Attacks, Intrusions and Defenses)

13. [arXiv:2607.28528](#) [[pdf](#)] [cs.CL](#)

AI systems and the reproduction of (standard) language ideologies in World Englishes

Authors: [Kingsley Ugwuanyi](#)

Abstract: The rapid growth of **large...** [▽ More](#)

Submitted 30 July, 2026; **originally announced** July 2026.

Comments: 13 pages, 0 figure

14. [arXiv:2607.24072](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

LLM-Based vs. Lexicon-Based Sentiment Signals for Tail-Risk Detection in Meme Stocks

Authors: [Paul Kilian](#), [Markus Kleffmann](#)

Abstract: This paper presents an empirical **comparison** of lexicon-based and... [▽ More](#)

Submitted 27 July, 2026; **originally announced** July 2026.

Comments: 10 pages, 1 figure, 4 tables

15. [arXiv:2607.11890](#) [[pdf](#)] [cs.CY](#) [cs.AI](#) [doi](#) [10.54536/ari.v4i1.6586](#)

So Many Opinions, So Many LLMs: **Comparing Large Language Models** to Traditional Machine Learning for Open- Ended Survey Analysis

Authors: [Abdullah Akinde](#), [Mariam Akinde](#), [Rasheedat Emiola](#), [Ahmed Akinsola](#)

Abstract: ...to classify text ("So Many Responses, So Little Time: A Machine-Learning Approach to Analyzing Open-Ended Survey Data") [1], this study investigates how different **large...** [▽ More](#)

Submitted 3 May, 2026; **originally announced** July 2026.

16. [arXiv:2607.11597](#) [[pdf](#)] [cs.CL](#)

Beyond **Benchmarks**: Exposing the Hidden Crisis in Bangla Hate Speech Detection

Authors: [Faria Afrin Tisha](#), [Fariya Tabassum](#), [Hafsa Binte Kibria](#), [Md. Nahiduzzaman](#), [Mominul Ahsan](#)

Abstract: The spread of hate speech (HS) across different **social...** [▽ More](#)

Submitted 13 July, 2026; **originally announced** July 2026.

17. [arXiv:2607.11279](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

FAD-SA-GRU: Enhancing Hate Speech Detection in Algerian Dialect Through Feature-Augmented Self-Attention GRU Networks

Authors: [Sara Yakoubi](#), [Ikram Khalfallah](#), [Kenza Khelkhal](#), [Dihia Lanasri](#)

Abstract: The widespread adoption of **social...** [▽ More](#)

Submitted 13 July, 2026; **originally announced** July 2026.

18. [arXiv:2607.06875](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.LG](#)

Video2Reaction: Mapping Video to Audience Reaction Distribution in the Wild

Authors: [Trang Nguyen](#), [Sidong Zhang](#), [Shiv Shankar](#), [Gauri Jagatap](#), [Deepak Chandran](#), [Andrea Fanelli](#), [Madalina Fiterau](#)

Abstract: ...of viewers in the wild, as expressed through **social media**. **Video2Reaction** spans more than 10,000 videos and serves as a reliable... [▽ More](#)

Submitted 7 July, 2026; **originally announced** July 2026.

19. [arXiv:2607.03981](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CV](#)

BanglaMemeEvidence: A Multimodal **Benchmark** Dataset for Explanatory Evidence Detection in Bengali Memes

Authors: [Fatema Tuj Johora Faria](#), [Mukaffi Bin Moin](#), [Md. Mahfuzur Rahman](#), [Pronay Debnath](#), [Asif Iftekher Fahim](#), [Faisal Muhammad Shah](#)

Abstract: Memes have become influential communication tools on **social...** [▽ More](#)

Submitted 4 July, 2026; **originally announced** July 2026.

Comments: Accepted at 6th International Conference on Innovations in Computational Intelligence and Computer Vision (ICICV 2026)

20. [arXiv:2606.30259](#) [[pdf](#)] [cs.CL](#)

Multi-Agent System Leveraging Open-Source LLMs to Mitigate Disinformation Threats

Authors: [Sebastian Kula](#), [Martin Tamajka](#)

Abstract: In contemporary societies, the threat of disinformation has reached alarming levels, exacerbated by the proliferation of electronic communication, **social...** [▽ More](#)

Submitted 29 June, 2026; **originally announced** June 2026.

21. [arXiv:2606.27314](#) [pdf, ps, other] cs.CL

Beyond Surface Forms: A Comprehensive, Mechanism-Oriented Taxonomy of Indirect Linguistic Encoding for LLM-Based Coded Language Detection

Authors: [Hamid Reza Firoozfar](#), [Mohammadsadegh Abolhasani](#), [Reza Mousavi](#), [Paul Jen-Hwa Hu](#)

Abstract: To avoid moderation and surveillance on **social**... [▽ More](#)

Submitted 13 September, 2026; v1 submitted 25 June, 2026; originally announced June 2026.

Comments: Submitted for review in ARR for EMNLP 2026

ACM Class: I.2.7

22. [arXiv:2606.27247](#) [pdf, ps, other] cs.LG

RSPC: A **Benchmark** for Modeling Stress and Psychiatric Conditions in Digitally Mediated Relationships using Psychiatrist **Annotations**

Authors: [Parmitha Vangapandu](#), [Sai Ganesh Mokkaapati](#), [Sathwik Narkedimilli](#), [MSVPJ Sathvik](#), [Timothy Liu](#), [Simon See](#), [Johannes C. Eichstaedt](#)

Abstract: In NLP, mental health conditions are often modeled as isolated phenomena, without interpersonal context. We use **Reddit** posts about long-distance relationships to capture both mental health distress and associated relational triggers. We introduce the Relational Stress and Psychiatry Corpus (RSPC) containing 1,799... [▽ More](#)

Submitted 25 June, 2026; originally announced June 2026.

23. [arXiv:2606.26205](#) [pdf, ps, other] cs.AI

Knowledge-augmented Agentic AI for Mental Health Medication Information Seeking

Authors: [Huizi Yu](#), [Jian Liu](#), [Wenkong Wang](#), [Lingyao Li](#), [Jiayan Zhou](#), [Zhaoqian Xue](#), [Xiang Li](#), [Xinxin Lin](#), [Zhiying Liang](#), [Zhuoru Wu](#), [Siyuan Ma](#), [Xin Ma](#), [Lizhou Fan](#)

Abstract: ...information can amplify fear, nocebo responses, and non-adherence. Here we develop a provenance-aware, knowledge-graph-based multi-agent framework unifying 466,525 **Reddit** posts, 60,782 WebMD reviews, and twenty years of U.S. FDA Adverse Event Reporting System records for nine antidepressants. A... [▽ More](#)

Submitted 24 June, 2026; originally announced June 2026.

24. [arXiv:2606.21996](#) [pdf, ps, other] cs.SI cs.CY

Cultural Targets, Structural Frames, Binding Morals: A Cross-Lingual Audit of Online Hate in Multicultural Singapore

Authors: [Emilio Ferrara](#)

Abstract: From a Singapore-centric 2025 Facebook, **Reddit**, and YouTube corpus (31.0M items; 1.76M comments mentioning eleven identity groups), we... [▽ More](#)

Submitted 20 June, 2026; originally announced June 2026.

25. [arXiv:2606.18285](#) [pdf, ps, other] cs.SI cs.CY

RELIANCE: Curating and Evaluating Reproductive Health Information on **Social Media**

Authors: [Vaibhav Balloli](#), [Laura Peyton Ellis](#), [Vishala Mishra](#), [Alice Chi](#), [Alex Peahl](#), [Elizabeth Bondi-Kelly](#)

Abstract: **Social**... [▽ More](#)

Submitted 10 June, 2026; originally announced June 2026.

26. [arXiv:2606.06526](#) [[pdf](#), [ps](#), [other](#)] [cs.AI](#) [cs.LG](#)

CrowdMath: A Dataset of Crowdsourced Mathematical Research Discussions

Authors: [Sherin Muckatira](#), [Jesse Geneson](#), [Slava Gerovitch](#), [Pavel Etingof](#), [Mikhail Gronas](#), [Anna Rumshisky](#)

Abstract: **Large...** [▽ More](#)

Submitted 28 August, 2026; v1 submitted 2 June, 2026; **originally announced** June 2026.

Comments: 16 pages, 4 figures; Accepted to EMNLP 2026 Findings

27. [arXiv:2606.06022](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Contextualized Prompting For **Stance Detection** On **Social Media**

Authors: [Tilman Beck](#), [Shakib Yazdani](#), [Simon Kruschinski](#), [Marcus Maurer](#), [Iryna Gurevych](#)

Abstract: **Stance...** [▽ More](#)

Submitted 4 June, 2026; **originally announced** June 2026.

28. [arXiv:2606.05256](#) [[pdf](#)] [cs.AI](#)

How Far Did They Go? The Persuasive Tactics of Covert LLM Agents in a Discontinued Field Experiment

Authors: [Kokil Jaidka](#), [Saifuddin Ahmed](#)

Abstract: This study analyzes a publicly released dataset from a discontinued field experiment on **Reddit's** r/ChangeMyView. The intervention, conducted by unknown, external researchers and halted following ethical backlash, involved undisclosed AI-generated accounts engaging users in live debate. After public disclosure,... [▽ More](#)

Submitted 3 June, 2026; **originally announced** June 2026.

29. [arXiv:2606.04867](#) [[pdf](#), [ps](#), [other](#)] [cs.AI](#)

AICompanionBench: **Benchmarking** LLMs-as-Judges for AI Companion Safety

Authors: [Yanjing Ren](#), [Reza Ebrahimi](#), [TengTeng Ma](#)

Abstract: ...Character.AI rapidly grow, concerns about unsafe human-AI interactions have intensified. This study introduces AICompanionBench, to our knowledge the first publicly available **benchmark** dataset of human-AI companion conversations... [▽ More](#)

Submitted 3 June, 2026; **originally announced** June 2026.

30. [arXiv:2606.00193](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

BOUTEF: A Multilingual Corpus for FakeNews in North Africa -- Language as a Weapon

Authors: [Kamel Smaili](#), [Yassine Toughrai](#), [Amina Laggoun](#), [David Langlois](#)

Abstract: The rapid spread of fake news on **social...** [▽ More](#)

Submitted 11 June, 2026; v1 submitted 29 May, 2026; **originally announced** June 2026.

31. [arXiv:2606.00046](#) [[pdf](#), [ps](#), [other](#)] [cs.MM](#) [cs.AI](#) [cs.CV](#) [cs.CY](#)

When Jokes Cross the Line: Analyzing Regular Humor and Dark Humor in YouTube Shorts

Authors: [Sydney Johns](#), [Sanjeev Parthasarathy](#), [Shantnu Bhalla](#), [Vaibhav Garg](#)

Abstract: ...negative effects on some audiences. To study this problem, we introduce TwistedHumor, a dataset of 1,211 YouTube Shorts paired with 33,041 related comments, with hand **annotations** for humor presence, humor type, harm, topic, rhetorical devices, and stand up context. Beyond dataset creation, we present a multi view analysis of how humor and harm appear in shor... [▽ More](#)

Submitted 30 April, 2026; originally announced June 2026.

32. [arXiv:2605.26445](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Curation and Extraction of Drug-Related Entities from **Reddit** Platform

Authors: [Zewei Wang](#), [Zihan Xu](#), [Yishu Wei](#), [Michael Chary](#), [Yifan Peng](#)

Abstract: ...of real-world usage. Meanwhile, drug users share first-hand experiences online, offering insights into dosage and effects of drugs. To bridge this gap, we introduce ReDose (**Reddit** Drug DOSe and Effect), a dataset of 6,435... [▽ More](#)

Submitted 8 April, 2026; originally announced May 2026.

Comments: Accepted by IEEE International Conference on Healthcare Informatics (ICHI 2026)

33. [arXiv:2605.25758](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

StreamProfileBench: A **Benchmark** for Fine-Grained User Profile Inference in Real-World Streaming Scenarios

Authors: [Sizhe Wang](#), [Feiyu Duan](#), [Juelin Wang](#), [Liwen Zhang](#), [Zhongyu Wei](#)

Abstract: **Large**... [▽ More](#)

Submitted 26 July, 2026; v1 submitted 25 May, 2026; originally announced May 2026.

34. [arXiv:2605.24675](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.AI](#) [doi](#) [10.1145/3770855.3817631](#)

VaaWIT: Visual-Aware Adaptation of **Large Language Models** for Multilingual Web Image Translation

Authors: [Bo Li](#), [Ronghao Chen](#), [Ningyuan Deng](#), [Huacan Wang](#), [Shaolin Zhu](#), [Lijie Wen](#)

Abstract: Translating text embedded in Web images is crucial for improving content accessibility and cross-lingual information retrieval, particularly within **social**... [▽ More](#)

Submitted 23 May, 2026; originally announced May 2026.

Comments: Accepted by KDD 2026

35. [arXiv:2605.23325](#) [[pdf](#)] [cs.CL](#)

AraHopeCorpus: **Annotation** Guidelines and Dataset for Hope Speech in Arabic **Social Media** Crisis Discourse

Authors: [Esra'a Sharqawi](#), [Wajdi Zaghouani](#)

Abstract: **Social**... [▽ More](#)

Submitted 22 May, 2026; originally announced May 2026.

36. [arXiv:2605.21097](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

WCXB: A Multi-Type Web Content Extraction **Benchmark**

Authors: [Murrough Foley](#)

Abstract: ...- isolating a page's main content from surrounding boilerplate - is a prerequisite for search indexing, retrieval-augmented generation, NLP dataset construction, and **large**... [▽ More](#)

Submitted 20 May, 2026; originally announced May 2026.

Comments: Dataset: github.com/Murrough-Foley/web-content-extraction-benchmark,
doi.org/10.5281/zenodo.19316874. Leaderboard: webcontentextraction.org. Preprint also deposited at
doi.org/10.5281/zenodo.19664685

37. [arXiv:2605.13663](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CY](#)

Fine-tuning with Hierarchical Prompting for Robust Propaganda Classification Across **Annotation** Schemas

Authors: [Lukas Stähelin](#), [Veronika Solopova](#), [Max Upravitelev](#), [David Kaplan](#), [Ariana Sahitaj](#), [Premtim Sahitaj](#), [Charlott Jakob](#), [Sebastian Möller](#), [Vera Schmitt](#)

Abstract: Propaganda detection in **social**... [▽ More](#)

Submitted 13 May, 2026; originally announced May 2026.

38. [arXiv:2605.12525](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.AI](#) [cs.CL](#)

PERCEIVE: A **Benchmark** for Personalized Emotion and Communication Behavior Understanding on **Social Media**

Authors: [Jian Liao](#), [Yujin Zheng](#), [Suge Wang](#), [Jianxing Zheng](#), [Deyu Li](#)

Abstract: Current emotion analysis in **social**... [▽ More](#)

Submitted 10 April, 2026; originally announced May 2026.

39. [arXiv:2605.12510](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.CL](#) [cs.CY](#)

WhatsApp Vaccine Discourse (WhaVax): An Expert-**Annotated** Dataset and **Benchmark** for Health Misinformation Detection

Authors: [Jônatas H. dos Santos](#), [Julio C. S. Reis](#), [Philipe Melo](#), [João F. H. Olivetti](#), [Thales H. Silva](#), [Matheus Gontijo Guimaraes](#), [Glaucio de Souza](#), [Marcos A. Gonçalves](#), [Fabricio Benevenuto](#), [Filipe B. B. Zanovello](#), [Marco A. G. Rodrigues](#), [Cristiano X. Lima](#)

Abstract: We introduce WhaVax, a new expert-**annotated** dataset of vaccine-related WhatsApp messages collected from large Brazilian public groups spanning multiple pandemic years. The dataset was constructed through a rigorous, carefully designed pipeline that integrates keyword-based data collection, semantic deduplication to remove near-duplicate content, and a multi-...

[▽ More](#)

Submitted 25 March, 2026; originally announced May 2026.

Comments: 10 pages. This is a preprint version of a paper accepted for the International AAAI Conference on Web and **Social Media** (ICWSM'26). Please cite the conference version rather than this preprint

40. [arXiv:2605.06940](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

MultiSoc-4D: A **Benchmark** for Diagnosing Instruction-Induced Label Collapse in Closed-Set LLM **Annotation** of Bengali **Social Media**

Authors: [Souvik Pramanik](#), [S. M. Riaz Rahman Antu](#), [Shak Mohammad Abyad](#), [Md. Ibrahim Khalil](#), [Md. Shahriar Hussain](#)

Abstract: **Annotation** automation via... [▽ More](#)

Submitted 7 May, 2026; originally announced May 2026.

Comments: 21 pages, 14 figures, 13 tables

41. [arXiv:2605.01323](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

SiNFluD: Creating and Evaluating Figurative Language Dataset for Sindhi

Authors: [Wazir Ali](#), [Adeeb Noor](#), [Saifullah Tumrani](#)

Abstract: In this article, we introduce SiNFluD, a novel **benchmark** dataset for Sindhi figurative language classification. We first collect raw text from various blogs,... [▽ More](#)

Submitted 9 May, 2026; v1 submitted 2 May, 2026; originally announced May 2026.

42. [arXiv:2605.00513](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

ControBench: An Interaction-Aware **Benchmark** for Controversial Discourse Analysis on Social Networks

Authors: [Ta Thanh Thuy](#), [Jiaqi Zhu](#), [Xuan Liu](#), [Lin Shang](#), [Reihaneh Rabbany](#), [Guillaume Rabusseau](#), [Lihui Chen](#), [Zheng Yilun](#), [Sitao Luan](#)

Abstract: ...structure, some model structure without rich semantics, and others represent conversations without stable user-level ideological identity. We introduce ControBench, a **benchmark** for controversial discourse analysis that combines heterogeneous social interaction graphs with rich textual semantics. Built from... [▽ More](#)

Submitted 1 May, 2026; originally announced May 2026.

43. [arXiv:2604.23458](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.IR](#) [cs.LG](#)

A **Benchmark** Suite of **Reddit**-Derived Datasets for Mental Health Detection

Authors: [Khalid Hasan](#), [Jamil Saquer](#)

Abstract: ...mental health through natural language processing (NLP). However, it is hindered by a lack of high-quality, well-validated datasets. Existing studies have a tendency to build **task**-... [▽ More](#)

Submitted 25 April, 2026; originally announced April 2026.

Comments: In the proceedings of 12th Annual Conference on Computational Science & Computational Intelligence (CSCI'25)

44. [arXiv:2604.21786](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#)

From Codebooks to VLMs: Evaluating Automated Visual Discourse Analysis for Climate Change on **Social Media**

Authors: [Katharina Prasse](#), [Steffen Jung](#), [Isaac Bravo](#), [Stefanie Walter](#), [Patrick Knab](#), [Christian Bartelt](#), [Margret Keuper](#)

Abstract: **Social**... [▽ More](#)

Submitted 23 April, 2026; originally announced April 2026.

45. [arXiv:2604.20511](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.AI](#) [cs.CL](#) [cs.CV](#) [cs.CY](#)

CHASM: Unveiling Covert Advertisements on Chinese **Social Media**

Authors: [Jingyi Zheng](#), [Tianyi Hu](#), [Yule Liu](#), [Zhen Sun](#), [Zongmin Zhang](#), [Zifan Peng](#), [Wenhan Dong](#), [Xinlei He](#)

Abstract: Current **benchmarks** for evaluating... [▽ More](#)

Submitted 22 April, 2026; originally announced April 2026.

Comments: NeuIPS 2025 (Datasets and **Benchmarks** Track)

46. [arXiv:2604.18955](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.SI](#)

Assessing Capabilities of **Large Language Models** in **Social Media** Analytics: A Multi-task Quest

Authors: [Ramtin Davoudi](#), [Kartik Thakkar](#), [Nazanin Donyapour](#), [Tyler Derr](#), [Hamid Karimi](#)

Abstract: ...we present the first comprehensive evaluation of modern LLMs - including GPT-4, GPT-4o, GPT-3.5-Turbo, Gemini 1.5 Pro, DeepSeek-V3, Llama 3.2, and BERT - across three core **social**...

[▽ More](#)

Submitted 20 April, 2026; originally announced April 2026.

47. [arXiv:2604.17319](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.CL](#)

E2E-GMNER: End-to-End Generative Grounded Multimodal Named Entity Recognition

Authors: [Meng Zhang](#), [Jinzhong Ning](#), [Xiaolong Wu](#), [Hongfei Lin](#), [Yijia Zhang](#)

Abstract: ...a fully end-to-end generative framework that unifies entity recognition, semantic typing, visual grounding, and implicit knowledge reasoning within a single multimodal **large...** [▽ More](#)

Submitted 19 April, 2026; originally announced April 2026.

Comments: Accepted to Findings of ACL 2026

48. [arXiv:2604.12813](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.MM](#)

DPC-VQA: Decoupling Quality Perception and Residual Calibration for Video Quality Assessment

Authors: [Xinyue Li](#), [Shubo Xu](#), [Zhichao Zhang](#), [Zhaolin Cai](#), [Yitong Chen](#), [Guangtao Zhai](#)

Abstract: Recent multimodal **large...** [▽ More](#)

Submitted 14 June, 2026; v1 submitted 14 April, 2026; originally announced April 2026.

49. [arXiv:2604.11121](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

BITS Pilani at SemEval-2026 Task 9: Structured Supervised Fine-Tuning with DPO Refinement for Polarization Detection

Authors: [Atharva Gupta](#), [Dhruv Kumar](#), [Yash Sinha](#)

Abstract: ...polarization. Accurate computational detection of online polarization is challenging due to nuanced rhetoric, implicit framing, and the high cost of human-in-the-loop **annotation**. Building on recent findings that contextual prompting enables... [▽ More](#)

Submitted 7 May, 2026; v1 submitted 13 April, 2026; originally announced April 2026.

Comments: Accepted to the 20th International Workshop on Semantic **Evaluation** (SemEval-2026), to be held in conjunction with ACL 2026

50. [arXiv:2604.09626](#) [[pdf](#), [ps](#), [other](#)] [cs.CY](#) [cs.LG](#) [cs.SI](#)

Investigating Vaccine Buyer's Remorse: Post-Vaccination Decision Regret in COVID-19 **Social Media** Using Politically Diverse Human **Annotation**

Authors: [Miles Stanley](#), [Soumyajit Datta](#), [Ashutosh Kumar](#), [Ashiqur R. KhudaBukhsh](#)

Abstract: ...and refining public health communication. In this paper, we curate a novel dataset from a large YouTube news corpus capturing COVID-19 vaccination experiences, and construct a **benchmark** subset focused on vaccine regret,... [▽ More](#)

Submitted 18 March, 2026; originally announced April 2026.

51. [arXiv:2604.00518](#) [[pdf](#), [ps](#), [other](#)] [cs.CY](#)

Do Agents Repair When Challenged -- or Just Reply? Challenge, Repair, and Public Correction in a Deployed Agent **Forum**

Authors: [Luyang Zhang](#), [Yi-Yun Chu](#), [Jialu Wang](#), [Beibei Li](#), [Ramayya Krishnan](#)

Abstract: As **large...** [▽ More](#)

Submitted 14 September, 2026; v1 submitted 1 April, 2026; originally announced April 2026.

52. [arXiv:2603.27011](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#)

ParsCN: A Persian Dataset for Counter-Narrative Generation to Combat Online Hate Speech

Authors: [Zahra Safdari Fesaghandis](#), [Suman Kalyan Maity](#)

Abstract: Online hate speech threatens online civility, particularly in **low-...** [▽ More](#)

Submitted 27 March, 2026; **originally announced** March 2026.

Comments: Accepted at the International AAAI Conference on Web and **Social Media** (ICWSM 2026); Paper Information: 16 pages, 3 Figures, 10 Tables

53. [arXiv:2603.23531](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Large Language Models Unpack Complex Political Opinions through Target-Stance Extraction

Authors: [Özgür Togay](#), [Javier Garcia-Bernardo](#), [Florian Kunneman](#), [Anastasia Giachanou](#)

Abstract: ...a wide range of subjects, making it difficult to automatically identify the target of discussion and the opinion expressed toward them. In this study, we investigate whether **Large...** [▽ More](#)

Submitted 20 May, 2026; v1 submitted 7 March, 2026; **originally announced** March 2026.

54. [arXiv:2603.21529](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.1145/3770854.3785698](#)

[doi](#) [10.1145/3770854.3785698](#) [doi](#) [10.1145/3770854.3785698](#)

SynSym: A Synthetic Data Generation Framework for Psychiatric Symptom Identification

Authors: [Migyeong Kang](#), [Jihyun Kim](#), [Hyolim Jeon](#), [Sunwoo Hwang](#), [Jihyun An](#), [Yonghoon Kim](#), [Haewoon Kwak](#), [Jisun An](#), [Jinyoung Han](#)

Abstract: Psychiatric symptom identification on **social...** [▽ More](#)

Submitted 22 March, 2026; **originally announced** March 2026.

ACM Class: I.2.7

55. [arXiv:2603.19744](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.1109/BigData66926.2025.11401919](#)

Rethinking Ground Truth: A Case Study on Human Label Variation in MLLM Benchmarking

Authors: [Tomas Ruiz](#), [Tanalp Agustoslu](#), [Carsten Schwemmer](#)

Abstract: Human Label Variation (HLV), i.e. systematic differences among **annotators'** judgments, remains underexplored in... [▽ More](#)

Submitted 20 March, 2026; **originally announced** March 2026.

Comments: 6 pages, 3 tables, 1 figure

Journal ref: 2025 IEEE International Conference on Big Data (BigData), 2025

56. [arXiv:2603.18611](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CV](#) [doi](#) [10.1145/3774904.3792991](#)

Cross-Modal Rationale Transfer for Explainable Humanitarian Classification on Social Media

Authors: [Thi Huyen Nguyen](#), [Koustav Rudra](#), [Wolfgang Nejdl](#)

Abstract: Advances in **social...** [▽ More](#)

Submitted 19 March, 2026; **originally announced** March 2026.

Comments: Accepted at WWW 2026

57. [arXiv:2603.14078](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

CMHL: Contrastive Multi-Head Learning for Emotionally Consistent **Text Classification**

Authors: [Menna Elgabry](#), [Ali Hamdi](#), [Khaled Shaban](#)

Abstract: Textual Emotion Classification (TEC) is one of the most difficult NLP tasks. State of the art approaches rely on **Large**... [▽ More](#)

Submitted 14 March, 2026; **originally announced** March 2026.

58. [arXiv:2603.12343](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

LLM-Augmented Therapy Normalization and Aspect-Based **Sentiment Analysis** for Treatment-Resistant Depression on **Reddit**

Authors: [Yuxin Zhu](#), [Sahithi Lakamana](#), [Masoud Rouhizadeh](#), [Selen Bozkurt](#), [Rachel Hershenberg](#), [Abeed Sarker](#)

Abstract: ...peer-support narratives therefore offer a complementary lens on how patients describe and evaluate medications in real-world use. In this study, we curated a corpus of 5,059 **Reddit** posts explicitly referencing TRD from 3,480 subscribers across 28 mental health-related subreddits from 2010 to 2025. Of these, 3,839 posts mentioned at least one medication, yiel... [▽ More](#)

Submitted 12 March, 2026; **originally announced** March 2026.

59. [arXiv:2602.21608](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

MixSarc: A Bangla-English Code-Mixed Corpus for Implicit Meaning Identification

Authors: [Kazi Samin Yasar Alam](#), [Md Tanbir Chowdhury](#), [Tamim Ahmed](#), [Ajwad Abrar](#), [Md Rafid Haque](#)

Abstract: Bangla-English code-mixing is widespread across South Asian **social**... [▽ More](#)

Submitted 27 June, 2026; **v1** submitted 25 February, 2026; **originally announced** February 2026.

Comments: Under Review

60. [arXiv:2602.19467](#) [[pdf](#)] [cs.CY](#) [cs.AI](#) [cs.CL](#)

Can **Large Language Models** Replace Human Coders? Introducing ContentBench

Authors: [Michael Haman](#)

Abstract: Can low-cost **large**... [▽ More](#)

Submitted 22 February, 2026; **originally announced** February 2026.

Comments: Project website: <https://contentbench.github.io>

61. [arXiv:2602.19333](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.IR](#) [cs.SI](#)

PerSoMed: A Large-Scale Balanced Dataset for Persian **Social Media Text Classification**

Authors: [Isun Chehreh](#), [Ebrahim Ansari](#)

Abstract: This research introduces the first large-scale, well-balanced Persian **social**... [▽ More](#)

Submitted 25 May, 2026; **v1** submitted 22 February, 2026; **originally announced** February 2026.

Comments: 10 pages, including 1 figure

ACM Class: I.2.7; H.3.1; H.3.3

62. [arXiv:2602.15842](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.CL](#)

Mememes-as-Replies: Can Models Select Humorous Manga Panel Responses?

Authors: [Ryosuke Kohita](#), [Seiichiro Yoshioka](#)

Abstract: ...use of memes to create humor remains an understudied area of web science. To address this gap, we introduce the Meme Reply Selection task and present MaMe-Re (Manga Meme Reply **Benchmark**), a... [▽ More](#)

Submitted 21 January, 2026; **originally announced** February 2026.

63. [arXiv:2602.14406](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

TruthStance: An **Annotated** Dataset of Conversations on Truth Social

Authors: [Fathima Ameen](#), [Danielle Brown](#), [Manusha Malgareddy](#), [Amanul Haque](#)

Abstract: Argument mining and **stance**... [▽ More](#)

Submitted 15 February, 2026; **originally announced** February 2026.

64. [arXiv:2602.13870](#) [[pdf](#)] [cs.CL](#)

ADAB: Arabic Dataset for Automated Politeness **Benchmarking** -- A Large-Scale Resource for Computational Sociopragmatics

Authors: [Hend Al-Khalifa](#), [Nadia Ghezaiel](#), [Maria Bounnit](#), [Hend Hamed Alhazmi](#), [Noof Abdullah Alfear](#), [Reem Fahad Alqifari](#), [Ameera Masoud Almasoud](#), [Sharefah Al-Ghamdi](#)

Abstract: ...under-explored, despite the rich and complex politeness expressions embedded in Arabic communication. In this paper, we introduce ADAB (Arabic Politeness Dataset), a new **annotated** Arabic dataset collected from four online platforms, including... [▽ More](#)

Submitted 21 February, 2026; **v1** submitted 14 February, 2026; **originally announced** February 2026.

Comments: Paper accepted @ The Fifteenth biennial Language Resources and **Evaluation** Conference (LREC2026)

65. [arXiv:2602.12301](#) [[pdf](#), [ps](#), [other](#)] [cs.SD](#) [cs.CL](#) [cs.IR](#) [cs.LG](#) [eess.AS](#)

Beyond Musical Descriptors: Extracting Preference-Bearing Intent in Music Queries

Authors: [Marion Baranes](#), [Romain Hennequin](#), [Elena V. Epure](#)

Abstract: Although **annotated** music descriptor datasets for user queries are increasingly common, few consider the user's intent behind these descriptors, which is essential for effectively meeting their needs. We introduce MusicRecoIntent, a manually... [▽ More](#)

Submitted 11 February, 2026; **originally announced** February 2026.

Comments: Accepted at NLP4MusA 2026 (4th Workshop on NLP for Music and Audio)

66. [arXiv:2602.11962](#) [[pdf](#), [ps](#), [other](#)] [cs.HC](#) [cs.CY](#)

Wisdom of the LLM Crowd: A Large Scale **Benchmark** of Multi-Label U.S. Election-Related Harmful **Social Media** Content

Authors: [Qile Wang](#), [Prerana Khatiwada](#), [Carolina Coimbra Vieira](#), [Benjamin E. Bagozzi](#), [Kenneth E. Barner](#), [Matthew Louis Mauriello](#)

Abstract: ...for understanding and potentially mitigating its downstream spread. In this study, we introduce USE24-XD, a large-scale dataset of nearly 100k posts collected from X (formerly **Twitter**) during the 2024 U.S. presidential election cycle, enriched with spatio-temporal metadata. To substantially reduce the cost of manual... [▽ More](#)

Submitted 23 February, 2026; **v1** submitted 12 February, 2026; **originally announced** February 2026.

67. [arXiv:2602.08371](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.18653/v1/2026.eacl-long.129](#)

ViGoEmotions: A **Benchmark** Dataset For Fine-grained Emotion Detection on Vietnamese Texts

Authors: [Hung Quang Tran](#), [Nam Tien Pham](#), [Son T. Luu](#), [Kiet Van Nguyen](#)

Abstract: Emotion classification plays a significant role in emotion prediction and harmful content detection. Recent advancements in NLP, particularly through **large...** [▽ More](#)

Submitted 26 March, 2026; **v1** submitted 9 February, 2026; **originally announced** February 2026.

Comments: Accepted as main paper at EACL 2026

68. [arXiv:2602.05189](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.HC](#) [cs.LG](#) [cs.SI](#)

Are Open-Weight LLMs Ready for **Social Media** Moderation? A **Comparative** Study on Bluesky

Authors: [Hsuan-Yu Chou](#), [Wajiha Naveed](#), [Shuyan Zhou](#), [Xiaowei Yang](#)

Abstract: As internet access expands, so does exposure to harmful content, increasing the need for effective moderation. Research has demonstrated that **large...** [▽ More](#)

Submitted 4 February, 2026; **originally announced** February 2026.

69. [arXiv:2601.19490](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.1007/978-3-032-21321-1_58](#)

ClaimPT: A Portuguese Dataset of **Annotated** Claims in News Articles

Authors: [Ricardo Campos](#), [Raquel Sequeira](#), [Sara Nerea](#), [Inês Cantante](#), [Diogo Folques](#), [Luís Filipe Cunha](#), [João Canavilhas](#), [António Branco](#), [Alípio Jorge](#), [Sérgio Nunes](#), [Nuno Guimarães](#), [Purificação Silvano](#)

Abstract: ...in automating fact-checking, where identifying claims is a crucial first step. However, progress has been uneven across languages, with English dominating due to abundant **annotated** data. Portuguese, like other languages, still lacks accessible, licensed datasets, limiting research, NLP developments and applications. In this paper, we introduce ClaimPT, a dat... [▽ More](#)

Submitted 26 March, 2026; **v1** submitted 27 January, 2026; **originally announced** January 2026.

Journal ref: Advances in Information Retrieval. ECIR 2026. Lecture Notes in Computer Science, vol 16486. Springer, Cham

70. [arXiv:2601.17601](#) [[pdf](#), [ps](#), [other](#)] [cs.IR](#) [cs.SI](#) [doi](#) [10.1145/3795766.3799753](#)

Why They Link: An Intent Taxonomy for Including Hyperlinks in Social Posts

Authors: [Fangping Lan](#), [Abdullah Aljebreen](#), [Eduard C. Dragut](#)

Abstract: URLs serve as bridges between **social...** [▽ More](#)

Submitted 2 March, 2026; **v1** submitted 24 January, 2026; **originally announced** January 2026.

Comments: 10 pages including references, 5 figures,

71. [arXiv:2601.08141](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

Qalb: Largest State-of-the-Art Urdu **Large Language Model** for 230M Speakers with Systematic Continued Pre-training

Authors: [Muhammad Taimoor Hassan](#), [Jawad Ahmed](#), [Muhammad Awais](#)

Abstract: Despite remarkable progress in **large...** [▽ More](#)

Submitted 12 January, 2026; **originally announced** January 2026.

72. [arXiv:2601.05232](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CY](#) [cs.LG](#)

AI Application Gives Users Real-Time Feedback on the Level of Peace in the **Social Media** Videos They Watch

Authors: P. Gilda, P. Dungarwal, A. Thongkham, E. T. Ajayi, S. Choudhary, T. M. Terol, C. Lam, J. P. Araujo, M. McFadyen-Mungalln, L. S. Liebovitch, P. T. Coleman, H. West, K. Sieck, S. Carter

Abstract: Most people now get their news from videos on **social media**, such as YouTube and Facebook, rather than through curated journalism. "We become what we behold." The content and tone of language plays an essential role in starting or ending conflicts. "Hate Speech" can enhance conflict, "Peace Speech" ... [▽ More](#)

Submitted 9 June, 2026; **v1** submitted 8 January, 2026; **originally announced** January 2026.

Comments: 6 pages, 4 figures, corrected typos, minor edits; v3: 16 pages, improved title, abstract, introduction, discussion, conclusions, added more references

ACM Class: I.2.6

73. [arXiv:2601.05091](#) [[pdf](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

Code-Mix **Sentiment Analysis** on Hinglish Tweets

Authors: Aashi Garg, Aneshya Das, Arshi Arya, Anushka Goyal, Aditi

Abstract: The effectiveness of brand monitoring in India is increasingly challenged by the rise of Hinglish--a hybrid of Hindi and English--used widely in **user**... [▽ More](#)

Submitted 8 January, 2026; **originally announced** January 2026.

Comments: Accepted at the 9th International Conference on Natural Language Processing and Information Retrieval (NLPPIR 2025), Fukuoka, Japan

74. [arXiv:2601.03194](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

X-MuTeST: A Multilingual **Benchmark** for Explainable Hate Speech Detection and A Novel LLM-consulted Explanation Framework

Authors: Mohammad Zia Ur Rehman, Sai Kartheek Reddy Kasu, Shashivardhan Reddy Koppula, Sai Rithwik Reddy Chirra, Shwetank Shekhar Singh, Nagendra Kumar

Abstract: Hate speech detection on **social**... [▽ More](#)

Submitted 6 January, 2026; **originally announced** January 2026.

Comments: Accepted in the proceedings of AAAI 2026

Journal ref: AAA 2026 (AIS)

75. [arXiv:2512.25015](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

MAMA-Memeia! Multi-Aspect Multi-Agent Collaboration for Depressive Symptoms Identification in Memes

Authors: Siddhant Agarwal, Adya Dhuler, Polly Ruhnke, Melvin Speisman, Md Shad Akhtar, Shweta Yadav

Abstract: ...the ever-growing utilization of memes in expressing depressive sentiments, we conduct a study on identifying depressive symptoms exhibited by memes shared by users of online **social**... [▽ More](#)

Submitted 31 December, 2025; **originally announced** December 2025.

Comments: Accepted by AAAI 2026

76. [arXiv:2512.21138](#) [[pdf](#)] [cs.SI](#)

Emotion Diffusion in Real and Simulated Social Graphs: Structural Limits of LLM-Based Social Simulation

Authors: Qiqi Qiang

Abstract: Understanding how emotions diffuse through social networks is central to computational social science. Recently, **large**... [▽ More](#)

Submitted 24 December, 2025; originally announced December 2025.

77. [arXiv:2512.19379](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.AI](#) [cs.MM](#)

OmniMER: Auxiliary-Enhanced LLM Adaptation for Indonesian Multimodal Emotion Recognition

Authors: [Xueming Yan](#), [Boyan Xu](#), [Yaochu Jin](#), [Lixian Xiao](#), [Wenlong Ye](#), [Runyang Cai](#), [Zeqi Zheng](#), [Jingfa Liu](#), [Aimin Yang](#), [Yongduan Song](#)

Abstract: Indonesian, spoken by over 200 million people, remains underserved in multimodal emotion recognition research despite its dominant presence on Southeast Asian **social**... [▽ More](#)

Submitted 10 February, 2026; v1 submitted 22 December, 2025; originally announced December 2025.

78. [arXiv:2512.14720](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.AI](#) [cs.CL](#)

SoMe: A Realistic **Benchmark** for LLM-based **Social Media** Agents

Authors: [Dizhan Xue](#), [Jing Cui](#), [Shengsheng Qian](#), [Chuanrui Hu](#), [Changsheng Xu](#)

Abstract: Intelligent agents powered by **large**... [▽ More](#)

Submitted 9 December, 2025; originally announced December 2025.

Comments: Accepted by AAAI 2026

79. [arXiv:2512.14565](#) [[pdf](#), [ps](#), [other](#)] [cs.IR](#)

Pairwise **Comparison** for Bias Identification and Quantification

Authors: [Fabian Haak](#), [Philipp Schaer](#)

Abstract: Linguistic bias in online news and **social**... [▽ More](#)

Submitted 16 December, 2025; originally announced December 2025.

MSC Class: 68T50 (Primary) 62J15 (Secondary) **ACM Class:** H.3.1; I.2.7

80. [arXiv:2512.00946](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Fine-tuning of lightweight **large language models** for **sentiment classification** on heterogeneous financial textual data

Authors: [Alvaro Paredes Amorin](#), [Andre Python](#), [Christoph Weisser](#)

Abstract: **Large**... [▽ More](#)

Submitted 30 November, 2025; originally announced December 2025.

81. [arXiv:2512.00673](#) [[pdf](#)] [cs.CL](#)

A **Comparison** of Human and ChatGPT Classification Performance on Complex **Social Media** Data

Authors: [Breanna E. Green](#), [Ashley L. Shea](#), [Pengfei Zhao](#), [Drew B. Margolin](#)

Abstract: ...among computational social scientists. Nevertheless, there remains space for improved understanding of the performance of ChatGPT in complex tasks such as classifying and **annotating** datasets containing nuanced language. Method. In this paper, we measure the performance of GPT-4 on one such task and... [▽ More](#)

Submitted 29 November, 2025; originally announced December 2025.

Comments: About 15 pages, draft version of accepted conference full paper. Published paper to follow

82. [arXiv:2511.23264](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.CL](#)

BanglaSentNet: An Explainable Hybrid Deep Learning Framework for Multi-Aspect **Sentiment Analysis** with Cross-Domain Transfer Learning

Authors: [Ariful Islam](#), [Md Rifat Hossen](#), [Tanvir Mahmud](#)

Abstract: Multi-aspect **sentiment...** [▽ More](#)

Submitted 28 November, 2025; **originally announced** November 2025.

Comments: Submitted to Springer Nature Computer Science (SNCS) as an extended version of our ICDSAIA 2025 conference paper

83. [arXiv:2511.22769](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Modeling Romanized Hindi and Bengali: Dataset Creation and Multilingual LLM Integration

Authors: [Kanchon Gharami](#), [Quazi Sarwar Muhtaseem](#), [Deepti Gupta](#), [Lavanya Elluri](#), [Shafika Showkat Moni](#)

Abstract: ...transliteration techniques to enhance the effectiveness of transforming Romanized scripts into native scripts is crucial for Natural Language Processing tasks, including **sentiment...** [▽ More](#)

Submitted 27 November, 2025; **originally announced** November 2025.

Comments: Proceedings of the 8th Workshop on Big Data for Cybersecurity (BigCyber)

84. [arXiv:2511.18146](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.1007/978-3-032-10209-6_11](#)

GeeSanBhava: Sentiment Tagged Sinhala Music Video Comment Data Set

Authors: [Yomal De Mel](#), [Nisansa de Silva](#)

Abstract: ...introduce GeeSanBhava, a high-quality data set of Sinhala song comments extracted from YouTube manually tagged using Russells Valence-Arousal model by three independent human **annotators**. The human... [▽ More](#)

Submitted 22 November, 2025; **originally announced** November 2025.

Journal ref: International Conference on Computational Collective Intelligence (2025)

85. [arXiv:2511.13759](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.AI](#)

Multi-Agent VLMs Guided Self-Training with PNU Loss for **Low-Resource** Offensive Content Detection

Authors: [Han Wang](#), [Deyi Ji](#), [Junyu Lu](#), [Lanyun Zhu](#), [Hailong Zhang](#), [Haiyang Wu](#), [Liqun Liu](#), [Peng Shu](#), [Roy Ka-Wei Lee](#)

Abstract: Accurate detection of offensive content on **social...** [▽ More](#)

Submitted 14 November, 2025; **originally announced** November 2025.

Comments: 8 pages, 4 figures, Fortieth AAAI Conference on Artificial Intelligence (AAAI-26)

ACM Class: I.2.10

86. [arXiv:2511.13238](#) [[pdf](#)] [cs.LG](#) [cs.AI](#) [cs.CL](#) [cs.CY](#) [doi](#) [10.1007/s11135-025-02500-4](#)

Computational Measurement of Political Positions: A Review of Text-Based Ideal Point Estimation Algorithms

Authors: [Patrick Parschan](#), [Charlott Jakob](#)

Abstract: ...in political science, communication, computational social science, and computer science to estimate ideological preferences from parliamentary speeches, party manifestos, and **social...** [▽ More](#)

Submitted 17 November, 2025; **originally announced** November 2025.

Comments: 46 pages, 8 figures, 2 tables, accepted for publication in Quality & Quantity

87. [arXiv:2511.03635](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

Towards Transparent **Stance Detection**: A **Zero-Shot** Approach Using Implicit and Explicit Interpretability

Authors: [Apoorva Upadhyaya](#), [Wolfgang Nejdl](#), [Marco Fisichella](#)

Abstract: **Zero**-... [▽ More](#)

Submitted 5 November, 2025; **originally announced** November 2025.

Comments: Accepted in AAAI CONFERENCE ON WEB AND **SOCIAL MEDIA** (ICWSM 2026)

88. [arXiv:2510.25273](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Adapting Small Language Models to **Low-Resource** Domains: A Case Study in Hindi Tourism QA

Authors: [Sandipan Majhi](#), [Paheli Bhattacharya](#)

Abstract: Domain-specific question answering in **low**-... [▽ More](#)

Submitted 29 October, 2025; **originally announced** October 2025.

Comments: Accepted at the **Forum** for Information Retrieval **Evaluation** 2025 (VATIKA Track)

89. [arXiv:2510.24178](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

MuSaG: A Multimodal German Sarcasm Dataset with Full-Modal **Annotations**

Authors: [Aaron Scott](#), [Maike Züfle](#), [Jan Niehues](#)

Abstract: Sarcasm is a complex form of figurative language in which the intended meaning contradicts the literal one. Its prevalence in **social**... [▽ More](#)

Submitted 4 March, 2026; **v1** submitted 28 October, 2025; **originally announced** October 2025.

90. [arXiv:2510.18582](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Beyond the Explicit: A Bilingual Dataset for Dehumanization Detection in **Social Media**

Authors: [Dennis Assenmacher](#), [Paloma Piot](#), [Katarina Laken](#), [David Jurgens](#), [Claudia Wagner](#)

Abstract: ...offensiveness, making them harder to detect yet equally damaging. Recognizing this gap, we use different sampling methods to collect a theory-informed bilingual dataset from **Twitter** and... [▽ More](#)

Submitted 21 October, 2025; **originally announced** October 2025.

91. [arXiv:2510.16985](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

[doi](#) [10.1109/COMPAS67506.2025.11381789](#)

Parameter-Efficient Fine-Tuning for **Low-Resource** Languages: A **Comparative** Study of LLMs for Bengali Hate Speech Detection

Authors: [Akif Islam](#), [Mohd Ruhul Ameen](#)

Abstract: Bengali **social**... [▽ More](#)

Submitted 19 October, 2025; **originally announced** October 2025.

Comments: Accepted to IEEE COMPAS 2025. 6 pages, 3 figures, 6 tables

Journal ref: 2025 IEEE International Conference on Computing, Applications and Systems (COMPAS), 2025

92. [arXiv:2510.11233](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

CNSocialDepress: A Chinese **Social Media** Dataset for Depression Risk Detection and Structured Analysis

Authors: Jinyuan Xu, Tian Lan, Xintao Yu, Xue He, Hezhi Zhang, Ying Wang, Pierre Magistry, Mathieu Valette, Lei Li

Abstract: ...Chinese-language resources for depression risk detection remain scarce and largely focus on binary classification. To address this limitation, we release CNSocialDepress, a **benchmark** dataset for depression risk detection on Chinese... [▼ More](#)

Submitted 11 May, 2026; **v1** submitted 13 October, 2025; **originally announced** October 2025.

93. [arXiv:2510.05291](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Camellia: **Benchmarking** Cultural Biases in LLMs for Asian Languages

Authors: Tarek Naous, Anagha Savit, Carlos Rafael Catalan, Geyang Guo, Jaehyeok Lee, Kyungdon Lee, Lheane Marie Dizon, Mengyu Ye, Neel Kothari, Sahajpreet Singh, Sarah Masud, Tanish Patwa, Trung Thanh Tran, Zohaib Khan, Alan Ritter, Tanmoy Chakraborty, Yuki Arase, Keisuke Sakaguchi, JinYeong Bak, Wei Xu

Abstract: As **Large**... [▼ More](#)

Submitted 31 August, 2026; **v1** submitted 6 October, 2025; **originally announced** October 2025.

Comments: Accepted at EMNLP 2026 Main Conference

94. [arXiv:2510.03502](#) [[pdf](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

ALHD: A Large-Scale and Multigenre **Benchmark** Dataset for Arabic LLM-Generated Text Detection

Authors: Ali Khairallah, Arkaitz Zubiaga

Abstract: We introduce ALHD, the first large-scale comprehensive Arabic dataset explicitly designed to distinguish between human- and LLM-generated texts. ALHD spans three genres (news, **social**... [▼ More](#)

Submitted 21 October, 2025; **v1** submitted 3 October, 2025; **originally announced** October 2025.

Comments: 47 pages, 15 figures. Dataset available at Zenodo: <https://doi.org/10.5281/zenodo.17249602> Codebase available at GitHub: <https://github.com/alikhairallah/ALHD-Benchmarking>

95. [arXiv:2510.01995](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

LLM-Based Multi-Task Bangla Hate Speech Detection: Type, Severity, and Target

Authors: Md Arid Hasan, Firoj Alam, Md Fahad Hossain, Usman Naseem, Syed Ishtiaque Ahmed

Abstract: Online **social**... [▼ More](#)

Submitted 2 October, 2025; **originally announced** October 2025.

MSC Class: 68T50 **ACM Class:** F.2.2; I.2.7

96. [arXiv:2509.25868](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.18653/v1/2026.eacl-long.381](https://doi.org/10.18653/v1/2026.eacl-long.381)

ReFACT: A **Benchmark** for Scientific Confabulation Detection with Positional Error **Annotations**

Authors: Yindong Wang, Martin Preiß, Margarita Bugueño, Jan Vincent Hoffbauer, Abdullatif Ghajar, Tolga Buz, Gerard de Melo

Abstract: The mechanisms underlying scientific confabulation in **Large**... [▼ More](#)

Submitted 23 April, 2026; **v1** submitted 30 September, 2025; **originally announced** September 2025.

Comments: Accepted to EACL 2026 (Main Conference, Oral presentation)

Journal ref: Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers), pages 8174-8187, Rabat, Morocco, 2026

97. [arXiv:2509.23515](#) [pdf, ps, other] cs.CL cs.AI

From Human Annotation to Automation: LLM-in-the-Loop Active Learning for Arabic Sentiment Analysis

Authors: [Dania Refai](#), [Alaa Dalaq](#), [Doaa Dalaq](#), [Irfan Ahmad](#)

Abstract: Natural language processing (NLP), particularly sentiment... [▽ More](#)

Submitted 27 September, 2025; originally announced September 2025.

98. [arXiv:2509.14922](#) [pdf] cs.CL

A Comparative Evaluation of Large Language Models for Persian Sentiment Analysis and Emotion Detection in Social Media Texts

Authors: [Kian Tohidi](#), [Kia Dashtipour](#), [Simone Rebora](#), [Sevda Pourfaramarz](#)

Abstract: This study presents a comprehensive comparative evaluation of four state-of-the-art...

[▽ More](#)

Submitted 18 September, 2025; originally announced September 2025.

Comments: 19 pages, 8 Figures, 9 Tables

99. [arXiv:2509.11802](#) [pdf] cs.CL

When Curiosity Signals Danger: Predicting Health Crises Through Online Medication Inquiries

Authors: [Dvora Goncharok](#), [Arbel Shifman](#), [Alexander Apartsin](#), [Yehudit Aperstein](#)

Abstract: Online medical forums are a rich and underutilized source of insight into patient concerns, especially regarding medication use. Some of the many questions users pose may signal confusion, misuse, or even the early warning signs of a developing health crisis. Detecting these critical questions that may precede severe adverse events or life-threatening compli... [▽ More](#)

Submitted 15 September, 2025; originally announced September 2025.

Comments: 5 pages, 2 figures

100. [arXiv:2509.11101](#) [pdf, ps, other] cs.CL

Seeing is Not Understanding: A Benchmark on Perception-Cognition Disparities in Large Language Models

Authors: [Haokun Li](#), [Yazhou Zhang](#), [Jizhi Ding](#), [Qiuchi Li](#), [Peng Zhang](#)

Abstract: With the rapid advancement of Multimodal Large... [▽ More](#)

Submitted 22 September, 2025; v1 submitted 14 September, 2025; originally announced September 2025.

Comments: I need to modify the content of the article

101. [arXiv:2509.10584](#) [pdf, ps, other] cs.CY cs.AI cs.CL

Smart Trial: Evaluating the Use of Large Language Models for Recruiting Clinical Trial Participants via Social Media

Authors: [Xiaofan Zhou](#), [Zisu Wang](#), [Janice Krieger](#), [Mohan Zalake](#), [Lu Cheng](#)

Abstract: ...time-consuming and geographically constrained. This work addresses the recruitment challenge by leveraging the vast amount of health-related information individuals share on social...

[▽ More](#)

Submitted 11 September, 2025; originally announced September 2025.

102. [arXiv:2509.10108](#) [pdf, ps, other] cs.CL

Scaling Arabic Medical Chatbots Using Synthetic Data: Enhancing Generative AI with Synthetic Patient Records

Authors: [Abdulrahman Allam](#), [Seif Ahmed](#), [Ali Hamdi](#), [Khaled Shaban](#)

Abstract: The development of medical chatbots in Arabic is significantly constrained by the scarcity of large-scale, high-quality **annotated** datasets. While prior efforts compiled a dataset of 20,000 Arabic patient-doctor interactions from... [▽ More](#)

Submitted 12 September, 2025; **originally announced** September 2025.

Comments: Accepted in AICCSA 2025

103. [arXiv:2509.09711](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

PsychiatryBench: A Multi-Task **Benchmark** for LLMs in Psychiatry

Authors: [Aya E. Fouda](#), [Abdelrahman A. Hassan](#), [Radwa J. Hanafy](#), [Mohammed E. Fouda](#)

Abstract: **Large**... [▽ More](#)

Submitted 23 November, 2025; **v1** submitted 7 September, 2025; **originally announced** September 2025.

104. [arXiv:2509.08960](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

BRoverbs -- Measuring how much LLMs understand Portuguese proverbs

Authors: [Thales Sales Almeida](#), [Giovana Kerche Bonás](#), [João Guilherme Alves Santos](#)

Abstract: **Large**... [▽ More](#)

Submitted 10 September, 2025; **originally announced** September 2025.

105. [arXiv:2509.08024](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.CY](#)

Two Stage Context Learning with **Large Language Models** for Multimodal **Stance Detection** on Climate Change

Authors: [Lata Pangtey](#), [Omkar Kabde](#), [Shahid Shafi Dar](#), [Nagendra Kumar](#)

Abstract: With the rapid proliferation of information across digital platforms, **stance**... [▽ More](#)

Submitted 9 September, 2025; **originally announced** September 2025.

106. [arXiv:2508.21675](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CV](#) [cs.GR](#)

Is this chart lying to me? Automating the detection of misleading visualizations

Authors: [Jonathan Tonglet](#), [Jan Zimny](#), [Tinne Tuytelaars](#), [Iryna Gurevych](#)

Abstract: Misleading visualizations are a potent driver of misinformation on **social**... [▽ More](#)

Submitted 17 April, 2026; **v1** submitted 29 August, 2025; **originally announced** August 2025.

Comments: Camera-ready version accepted at ACL 2026 Main conference. Code and data available at: <https://github.com/UKPLab/acl2026-misviz>

107. [arXiv:2508.20148](#) [[pdf](#)] [cs.AI](#) [cs.HC](#) [cs.MA](#)

The Anatomy of a Personal Health Agent

Authors: [A. Ali Heydari](#), [Ken Gu](#), [Vidya Srinivas](#), [Hong Yu](#), [Zhihan Zhang](#), [Yuwei Zhang](#), [Akshay Paruchuri](#), [Qian He](#), [Hamid Palangi](#), [Nova Hammerquist](#), [Ahmed A. Metwally](#), [Brent Winslow](#), [Yubin Kim](#), [Kumar Ayush](#), [Yuzhe Yang](#), [Girish Narayanswamy](#), [Maxwell A. Xu](#), [Jake Garrison](#), [Amy Armento Lee](#), [Jenny Vafeiadou](#), [Ben Graef](#), [Isaac R. Galatzer-Levy](#), [Erik Schenck](#), [Andrew Barakat](#), [Javier Perez](#), et al. (13 additional authors not shown)

Abstract: Health is a fundamental pillar of human wellness, and the rapid advancements in **large**... [▽ More](#)

Submitted 18 September, 2025; v1 submitted 27 August, 2025; originally announced August 2025.

Comments: Minor updates to the manuscript (V2)

108. [arXiv:2508.19467](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.IR](#)

Inference Gap in Domain Expertise and Machine Intelligence in Named Entity Recognition: Creation of and Insights from a Substance Use-related Dataset

Authors: [Sumon Kanti Dey](#), [Jeanne M. Powell](#), [Azra Ismail](#), [Jeanmarie Perrone](#), [Abeed Sarker](#)

Abstract: Nonmedical opioid use is an urgent public health challenge, with far-reaching clinical and social consequences that are often underreported in traditional healthcare settings. [Social...](#) [▽ More](#)

Submitted 26 August, 2025; originally announced August 2025.

Comments: Dataset and code: https://github.com/SumonKantiDey/Reddit_Impacts_NER

109. [arXiv:2508.15825](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [q-fin.ST](#)

Enhancing Cryptocurrency [Sentiment Analysis](#) with Multimodal Features

Authors: [Chenghao Liu](#), [Aniket Mahanti](#), [Ranesh Naha](#), [Guanghao Wang](#), [Erwann Sbai](#)

Abstract: As cryptocurrencies gain popularity, the digital asset marketplace becomes increasingly significant. Understanding [social...](#) [▽ More](#)

Submitted 31 August, 2025; v1 submitted 17 August, 2025; originally announced August 2025.

110. [arXiv:2508.13804](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.HC](#) [doi](#) [10.18653/v1/2025.uncertainlp-main.3](#)

Beyond Human Judgment: A Bayesian Evaluation of LLMs' Moral Values Understanding

Authors: [Maciej Skorski](#), [Alina Landowska](#)

Abstract: How do [Large...](#) [▽ More](#)

Submitted 20 November, 2025; v1 submitted 19 August, 2025; originally announced August 2025.

Comments: Appears in UncertainNLP@EMNLP 2025

MSC Class: 68T50; 62F15; 62P25 ACM Class: I.2.7; K.4.1; J.4

111. [arXiv:2508.13196](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.AI](#) [cs.IR](#)

Contextual Attention-Based Multimodal Fusion of LLM and CNN for [Sentiment Analysis](#)

Authors: [Meriem Zerkouk](#), [Miloud Mihoubi](#), [Belkacem Chikhaoui](#)

Abstract: This paper introduces a novel approach for multimodal [sentiment...](#) [▽ More](#)

Submitted 15 August, 2025; originally announced August 2025.

Comments: The 38th Canadian Conference on Artificial Intelligence (2025)

112. [arXiv:2508.12605](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#)

ViDA-UGC: Detailed Image Quality Analysis via Visual Distortion Assessment for UGC Images

Authors: [Wenjie Liao](#), [Jieyu Yuan](#), [Yifang Xu](#), [Chunle Guo](#), [Zilong Zhang](#), [Jihong Li](#), [Jiachen Fu](#), [Haotian Fan](#), [Tao Li](#), [Junhui Cui](#), [Chongyi Li](#)

Abstract: Recent advances in Multimodal [Large...](#) [▽ More](#)

Submitted 18 August, 2025; originally announced August 2025.

113. [arXiv:2508.11715](#) [[pdf](#), [ps](#), [other](#)] [cs.SE](#) [cs.AI](#)

Benchmark Dataset Generation and Evaluation for Excel Formula Repair with LLMs

Authors: [Ananya Singha](#), [Harshita Sahijwani](#), [Walt Williams](#), [Emmanuel Aboah Boateng](#), [Nick Hausman](#), [Miguel Di Luca](#), [Keegan Choudhury](#), [Chaya Binet](#), [Vu Le](#), [Tianwei Chen](#), [Oryan Rokeah Chen](#), [Sulaiman Vesal](#), [Sadid Hasan](#)

Abstract: ...yet often complex tool, particularly for novice users, where runtime errors arising from logical mistakes or misinterpretations of functions pose a significant challenge. While **large**... [▽ More](#)
Submitted 14 August, 2025; **originally announced** August 2025.

Comments: Accepted at the KDD workshop on **Evaluation** and Trustworthiness of Agentic and Generative AI Models

114. [arXiv:2508.11166](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Overcoming **Low-Resource** Barriers in Tulu: Neural Models and Corpus Creation for Offensive Language Identification

Authors: [Anusha M D](#), [Deepthi Vikram](#), [Bharathi Raja Chakravarthi](#), [Parameshwar R Hegde](#)

Abstract: Tulu, a **low**-... [▽ More](#)

Submitted 14 August, 2025; **originally announced** August 2025.

Comments: 20 pages, 3 tables, 3 figures. Submitted to Language Resources and **Evaluation** (Springer)

115. [arXiv:2508.10001](#) [[pdf](#)] [cs.CL](#) [cs.AI](#)

HiFACTMix: A Code-Mixed **Benchmark** and Graph-Aware Model for EvidenceBased Political Claim Verification in Hinglish

Authors: [Rakesh Thakur](#), [Sneha Sharma](#), [Gauri Chopra](#)

Abstract: Fact-checking in code-mixed, **low**-... [▽ More](#)

Submitted 4 August, 2025; **originally announced** August 2025.

116. [arXiv:2508.08649](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [doi](#) [10.18653/v1/2024.wassa-1.6](#)

LLaMA-Based Models for Aspect-Based **Sentiment Analysis**

Authors: [Jakub Šmíd](#), [Pavel Přibáň](#), [Pavel Král](#)

Abstract: While **large**... [▽ More](#)

Submitted 12 August, 2025; **originally announced** August 2025.

Comments: Published in Proceedings of the 14th Workshop on Computational Approaches to Subjectivity, Sentiment, & **Social Media** Analysis (WASSA 2024). Official version: <https://aclanthology.org/2024.wassa-1.6/>

117. [arXiv:2508.06515](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#)

BigTokDetect: A Clinically-Informed Vision-Language Modeling Framework for Detecting Pro-Bigorexia Videos on TikTok

Authors: [Minh Duc Chu](#), [Kshitij Pawar](#), [Zihao He](#), [Roxanna Sharifi](#), [Ross Sonnenblick](#), [Magdalayna Curry](#), [Laura D'Adamo](#), [Lindsay Young](#), [Stuart B Murray](#), [Kristina Lerman](#)

Abstract: **Social**... [▽ More](#)

Submitted 24 January, 2026; **v1** submitted 30 July, 2025; **originally announced** August 2025.

118. [arXiv:2508.06435](#) [[pdf](#)] [cs.CL](#) [cs.AI](#)

Learning the Topic, Not the Language: How LLMs Classify Online Immigration Discourse Across Languages

Authors: [Andrea Nasuto](#), [Stefano Maria Iacus](#), [Francisco Rowe](#), [Devika Jain](#)

Abstract: **Large**... [▽ More](#)

119. [arXiv:2507.20924](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CY](#) [cs.SI](#)

FHSTP@EXIST 2025 **Benchmark**: Sexism Detection with Transparent Speech Concept Bottleneck Models

Authors: [Roberto Labadie-Tamayo](#), [Adrian Jaques Böck](#), [Djordje Slijepčević](#), [Xihui Chen](#), [Andreas Babic](#), [Matthias Zeppelzauer](#)

Abstract: Sexism has become widespread on **social**... [▽ More](#)

Submitted 28 July, 2025; originally announced July 2025.

Comments: 12 pages

ACM Class: I.2

120. [arXiv:2507.16331](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Re:Form -- Reducing Human **Annotations** in Scalable Formal Software Verification with RL in LLMs: A Preliminary Study on Dafny

Authors: [Chuanhao Yan](#), [Fengdi Che](#), [Xuhan Huang](#), [Xu Xu](#), [Xin Li](#), [Yizhi Li](#), [Xingwei Qu](#), [Jingzhe Shi](#), [Chenghua Lin](#), [Yaodong Yang](#), [Binhang Yuan](#), [Hang Zhao](#), [Yu Qiao](#), [Bowen Zhou](#), [Jie Fu](#)

Abstract: Existing informal language-based (e.g., human language) **Large**... [▽ More](#)

Submitted 3 July, 2026; v1 submitted 22 July, 2025; originally announced July 2025.

Comments: Published in Transactions on Machine Learning Research (TMLR), 05/2026. Reviewed on OpenReview: <https://openreview.net/forum?id=cAQmIS4GOe>

121. [arXiv:2507.14693](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CY](#) [cs.LG](#)

Rethinking Suicidal Ideation Detection: A Trustworthy **Annotation** Framework and Cross-Lingual Model Evaluation

Authors: [Amina Dzafic](#), [Merve Kavut](#), [Ulya Bayram](#)

Abstract: Suicidal ideation detection is critical for real-time suicide prevention, yet its progress faces two under-explored challenges: limited language coverage and unreliable **annotation** practices. Most available datasets are in English, but even among these, high-quality, human-... [▽ More](#)

Submitted 19 July, 2025; originally announced July 2025.

Comments: This manuscript has been submitted to the IEEE Journal of Biomedical and Health Informatics

122. [arXiv:2507.14681](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Large Language Models as Medical Codes Selectors: a **benchmark** using the International Classification of Primary Care

Authors: [Vinicius Anjos de Almeida](#), [Vinicius de Camargo](#), [Raquel Gómez-Bravo](#), [Egbert van der Haring](#), [Kees van Boven](#), [Marcelo Finger](#), [Luis Fernandez Lopez](#)

Abstract: Background: Medical coding structures healthcare data for research, quality monitoring, and policy. This study assesses the potential of **large**... [▽ More](#)

Submitted 1 November, 2025; v1 submitted 19 July, 2025; originally announced July 2025.

Comments: Accepted at NeurIPS 2025 as a poster presentation in The Second Workshop on GenAI for Health: Potential, Trust, and Policy Compliance (<https://openreview.net/forum?id=KI7KZwjFEG>). 33 pages, 10 figures (including appendix), 15 tables (including appendix). To be submitted to peer-reviewed journal. For associated code repository, see <https://github.com/almeidava93/llm-as-code-selectors-paper>

123. [arXiv:2507.11559](#) [[pdf](#), [ps](#), [other](#)] [cs.CY](#) [cs.SI](#)

RSD-15K: A Large-Scale User-Level **Annotated** Dataset for Suicide Risk Detection on **Social Media**

Authors: [Shouwen Zheng](#), [Yingzhi Tao](#), [Taiqi Zhou](#)

Abstract: ...especially the suicide problem caused by multiple factors such as social competition, economic pressure and interpersonal relationships among young and middle-aged people. **Social...**

[▽ More](#)

Submitted 14 July, 2025; originally announced July 2025.

Comments: the article has already been recieved by 2025 IEEE 41st International Conference on Data Engineering Workshops (ICDEW), but hadn't been online yet

124. [arXiv:2507.09485](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Balanced Training Data Augmentation for Aspect-Based **Sentiment Analysis**

Authors: [Junjie Liu](#), [Yuanhe Tian](#), [Yan Song](#)

Abstract: Aspect-based **sentiment...** [▽ More](#)

Submitted 13 July, 2025; originally announced July 2025.

125. [arXiv:2507.08169](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [doi](#) [10.32473/flairs.38.1.138706](#)

Analysis of Propaganda in Tweets From Politically Biased Sources

Authors: [Vivek Sharma](#), [Mohammad Mahdi Shokri](#), [Sarah Ita Levitan](#), [Elena Filatova](#), [Shweta Jain](#)

Abstract: ...with popular news outlets, in propagating their bias using some form of propaganda-like language. We introduce JMBX(Journalist Media Bias on X), a systematically collected and **annotated** dataset of 1874 tweets from... [▽ More](#)

Submitted 10 July, 2025; originally announced July 2025.

Comments: The International FLAIRS Conference Proceedings, 38(1). <https://doi.org/10.32473/flairs.38.1.138706>

126. [arXiv:2507.08104](#) [[pdf](#), [ps](#), [other](#)] [cs.MM](#) [cs.AI](#) [cs.CL](#) [cs.CV](#)

VideoConviction: A Multimodal **Benchmark** for Human Conviction and Stock Market Recommendations

Authors: [Michael Galarnyk](#), [Veer Kejriwal](#), [Agam Shah](#), [Yash Bhardwaj](#), [Nicholas Meyer](#), [Anand Krishnan](#), [Sudheer Chava](#)

Abstract: **Social...** [▽ More](#)

Submitted 4 June, 2025; originally announced July 2025.

127. [arXiv:2506.23524](#) [[pdf](#)] [cs.CL](#) [cs.AI](#)

NEU-ESC: A Comprehensive Vietnamese dataset for Educational **Sentiment analysis** and topic Classification toward multitask learning

Authors: [Phan Quoc Hung Mai](#), [Quang Hung Nguyen](#), [Phuong Giang Duong](#), [Hong Hanh Nguyen](#), [Nguyen Tuan Long](#)

Abstract: ...limited. Existing educational datasets often lack domain relevance and student slang. To address these gaps, we introduce NEU-ESC, a new Vietnamese dataset for Educational **Sentiment...**

[▽ More](#)

Submitted 30 June, 2025; originally announced June 2025.

128. [arXiv:2506.17693](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

Zero-Shot Conversational **Stance Detection**: Dataset and Approaches

Authors: [Yuzhe Ding](#), [Kang He](#), [Bobo Li](#), [Li Zheng](#), [Haijun He](#), [Fei Li](#), [Chong Teng](#), [Donghong Ji](#)

Abstract: **Stance...** [▽ More](#)

Submitted 21 June, 2025; originally announced June 2025.

Comments: ACL 2025 (Findings)

129. [arXiv:2506.13470](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Induce, Align, Predict: **Zero-Shot Stance Detection** via Cognitive Inductive Reasoning

Authors: [Bowen Zhang](#), [Jun Ma](#), [Fuqiang Niu](#), [Li Dong](#), [Jinzhou Cao](#), [Genan Dai](#)

Abstract: **Zero-**... [▽ More](#)

Submitted 26 January, 2026; v1 submitted 16 June, 2025; originally announced June 2025.

Comments: Accepted at AAAI 2026

MSC Class: I.2.7; I.2.6

130. [arXiv:2506.12077](#) [[pdf](#), [ps](#), [other](#)] [cs.CY](#) [cs.CL](#)

Artificial Intelligence and Civil Discourse: How LLMs Moderate Climate Change Conversations

Authors: [Wenlu Fan](#), [Wentao Xu](#)

Abstract: As **large**... [▽ More](#)

Submitted 6 June, 2025; originally announced June 2025.

Comments: 10 pages

131. [arXiv:2506.08584](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

CounselBench: A Large-Scale Expert Evaluation and Adversarial **Benchmarking of Large Language Models** in Mental Health Question Answering

Authors: [Yahan Li](#), [Jifan Yao](#), [John Bosco S. Bunyi](#), [Adam C. Frank](#), [Angel Hsing-Chi Hwang](#), [Ruishan Liu](#)

Abstract: Medical question answering (QA) **benchmarks** often focus on multiple-choice or fact-based tasks, leaving open-ended answers to real patient questions underexplored. This gap is particularly critical in mental health, where patient questions often mix symptoms, treatment concerns, and emotional needs, requiring answers that balance clinical caution with context...

[▽ More](#)

Submitted 14 May, 2026; v1 submitted 10 June, 2025; originally announced June 2025.

132. [arXiv:2506.08147](#) [[pdf](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

Multilingual Hate Speech Detection in **Social Media** Using Translation-Based Approaches with **Large Language Models**

Authors: [Muhammad Usman](#), [Muhammad Ahmad](#), [M. Shahiki Tash](#), [Irina Gelbukh](#), [Rolando Quintero Tellez](#), [Grigori Sidorov](#)

Abstract: **Social**... [▽ More](#)

Submitted 9 June, 2025; originally announced June 2025.

133. [arXiv:2506.04156](#) [[pdf](#)] [cs.CL](#) [doi](#) [10.1038/s41597-026-06639-z](#)

A Dataset for Addressing Patient's Information Needs related to Clinical Course of Hospitalization

Authors: [Sarvesh Soni](#), [Dina Demner-Fushman](#)

Abstract: ...relevance of AI-generated responses. To our knowledge, no existing dataset captures patient information needs in the context of their EHRs. We introduce ArchEHR-QA, an expert-**annotated** dataset based on real-world patient cases from intensive care unit and emergency department settings. The cases comprise questions posed by patients to public health... [▽ More](#)
Submitted 4 June, 2025; **originally announced** June 2025.

134. [arXiv:2505.18916](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

SCRum-9: Multilingual Stance Classification over Rumours on **Social Media**

Authors: [Yue Li](#), [Jake Vasilakes](#), [Zhixue Zhao](#), [Carolina Scarton](#)

Abstract: ...SCRum-9 goes beyond existing stance classification datasets by covering more languages, linking examples to more fact-checked claims (2.1k), and including confidence-related **annotations** from multiple... [▽ More](#)

Submitted 15 November, 2025; **v1** submitted 24 May, 2025; **originally announced** May 2025.

Comments: Accepted by ICWSM 2026

135. [arXiv:2505.18331](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

PerMedCQA: **Benchmarking Large Language Models** on Medical Consumer Question Answering in Persian Language

Authors: [Naghmeh Jamali](#), [Milad Mohammadi](#), [Danial Baledi](#), [Zahra Rezvani](#), [Hesham Faili](#)

Abstract: Medical consumer question answering (CQA) is crucial for empowering patients by providing personalized and reliable health information. Despite recent advances in **large**... [▽ More](#)

Submitted 23 May, 2025; **originally announced** May 2025.

136. [arXiv:2505.16455](#) [[pdf](#), [other](#)] [cs.AI](#) [cs.CY](#)

Psychology-driven LLM Agents for Explainable Panic Prediction on **Social Media** during Sudden Disaster Events

Authors: [Mengzhu Liu](#), [Zhengqiu Zhu](#), [Chuan Ai](#), [Chen Gao](#), [Xinghong Li](#), [Lingnan He](#), [Kaisheng Lai](#), [Yingfeng Chen](#), [Xin Lu](#), [Yong Li](#), [Quanjun Yin](#)

Abstract: During sudden disaster events, accurately predicting public panic sentiment on **social**... [▽ More](#)

Submitted 22 May, 2025; **originally announced** May 2025.

137. [arXiv:2505.15801](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

VerifyBench: **Benchmarking** Reference-based Reward Systems for **Large Language Models**

Authors: [Yuchen Yan](#), [Jin Jiang](#), [Zhenbang Ren](#), [Yijun Li](#), [Xudong Cai](#), [Yang Liu](#), [Xin Xu](#), [Mengdi Zhang](#), [Jian Shao](#), [Yongliang Shen](#), [Jun Xiao](#), [Yueting Zhuang](#)

Abstract: ...of reference-based reward systems within reinforcement learning (RL), where model outputs are evaluated against ground truth references. However, existing reward **benchmarks** focus on preference... [▽ More](#)

Submitted 18 February, 2026; **v1** submitted 21 May, 2025; **originally announced** May 2025.

Comments: ICLR 2026: <https://openreview.net/forum?id=JfsjGmuFxz> Project Page: <https://zju-real.github.io/VerifyBench> Dataset: <https://huggingface.co/datasets/ZJU-REAL/VerifyBench> Code: <https://github.com/ZJU-REAL/VerifyBench>

138. [arXiv:2505.15616](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#)

LENS: Multi-level Evaluation of Multimodal Reasoning with Large Language Models

Authors: Ruilin Yao, Bo Zhang, Jirui Huang, Xinwei Long, Yifang Zhang, Tianyu Zou, Yufei Wu, Shichao Su, Yifan Xu, Wenxi Zeng, Zhaoyu Yang, Guoyou Li, Shilan Zhang, Zichan Li, Yaxiong Chen, Shengwu Xiong, Peng Xu, Jiajun Zhang, Bowen Zhou, David Clifton, Luc Van Gool

Abstract: Multimodal **Large**... ▾ More

Submitted 13 May, 2026; **v1** submitted 21 May, 2025; **originally announced** May 2025.

Comments: Published as a conference paper at ICLR 2026

139. [arXiv:2505.15101](#) [[pdf](#), [ps](#), [other](#)] [cs.LG](#) [cs.CL](#) [cs.IT](#)

Cost-aware LLM-based Online Dataset Annotation

Authors: Eray Can Elumar, Cem Tekin, Osman Yagan

Abstract: Recent advances in **large**... ▾ More

Submitted 15 December, 2025; **v1** submitted 21 May, 2025; **originally announced** May 2025.

Journal ref: In The Thirty-ninth Annual Conference on Neural Information Processing Systems, 2025. URL <https://openreview.net/forum?id=3AdTRYA2uj>

140. [arXiv:2505.15071](#) [[pdf](#), [other](#)] [cs.CL](#)

Can Large Language Models Understand Internet Buzzwords Through User-Generated Content

Authors: Chen Huang, Junkai Luo, Xinzuo Wang, Wenqiang Lei, Jiancheng Lv

Abstract: The massive **user**-... ▾ More

Submitted 20 May, 2025; **originally announced** May 2025.

Comments: ACL 2025 Main Paper. Our dataset and code are available at <https://github.com/SCUNLP/Buzzword>

141. [arXiv:2505.14422](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#)

MindVote: When AI Meets the Wild West of Social Media Opinion

Authors: Xutao Mao, Ezra Xuanru Tao, Leyao Wang

Abstract: **Large**... ▾ More

Submitted 7 November, 2025; **v1** submitted 20 May, 2025; **originally announced** May 2025.

Comments: Accepted in AAAI 2026 (Oral)

142. [arXiv:2505.14165](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#)

PL-FGSA: A Prompt Learning Framework for Fine-Grained Sentiment Analysis Based on MindSpore

Authors: Zhenkai Qin, Jiajing He, Qiao Fang

Abstract: Fine-grained **sentiment**... ▾ More

Submitted 20 May, 2025; **originally announced** May 2025.

143. [arXiv:2505.14106](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

A Personalized Conversational Benchmark: Towards Simulating Personalized Conversations

Authors: Li Li, Peilin Cai, Ryan A. Rossi, Franck Dernoncourt, Branislav Kveton, Junda Wu, Tong Yu, Linxin Song, Tiankai Yang, Yuehan Qin, Nesreen K. Ahmed, Samyadeep Basu, Subhojyoti Mukherjee, Ruiyi Zhang, Zhengmian Hu, Bo Ni, Yuxiao Zhou, Zichao Wang, Yue Huang, Yu Wang, Xiangliang Zhang, Philip S. Yu, Xiyang Hu, Yue Zhao

Abstract: We present PersonaConvBench, a large-scale **benchmark** for evaluating personalized reasoning and generation in multi-turn conversations with... [▽ More](#)

Submitted 25 May, 2025; v1 submitted 20 May, 2025; originally announced May 2025.

144. [arXiv:2505.12116](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

A Multi-Task **Benchmark** for Abusive Language Detection in **Low-Resource** Settings

Authors: [Fitsum Gaim](#), [Hoyun Song](#), [Huije Lee](#), [Changgeon Ko](#), [Eui Jun Hwang](#), [Jong C. Park](#)

Abstract: ...the majority of the world's languages due to the lack of resources, leaving millions of vulnerable users to online hostility. This work presents a large-scale human-**annotated** multi-task...

[▽ More](#)

Submitted 25 October, 2025; v1 submitted 17 May, 2025; originally announced May 2025.

Comments: Accepted at NeurIPS 2025

ACM Class: I.2.7

145. [arXiv:2505.10260](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Comparing LLM Text **Annotation** Skills: A Study on Human Rights Violations in **Social Media** Data

Authors: [Poli Apollinaire Nemkova](#), [Solomon Ubani](#), [Mark V. Albert](#)

Abstract: In the era of increasingly sophisticated natural language processing (NLP) systems, **large**... [▽ More](#)

Submitted 15 May, 2025; originally announced May 2025.

146. [arXiv:2505.09665](#) [[pdf](#), [ps](#), [other](#)] [cs.SI](#) [cs.CL](#)

Tales of the 2025 Los Angeles Fire: Hotwash for Public Health Concerns in **Reddit** via LLM-Enhanced Topic Modeling

Authors: [Sulong Zhou](#), [Qunying Huang](#), [Shaoheng Zhou](#), [Yun Hang](#), [Xinyue Ye](#), [Aodong Mei](#), [Kathryn Phung](#), [Yuning Ye](#), [Uma Govindswamy](#), [Zehan Li](#)

Abstract: ...and severe in recent years. Understanding how affected populations perceive and respond during wildfire crises is critical for timely and empathetic disaster response. **Social**...

[▽ More](#)

Submitted 5 January, 2026; v1 submitted 14 May, 2025; originally announced May 2025.

Comments: Fix typos in Method Section. Add data/code availability

147. [arXiv:2505.08464](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.LG](#) [cs.SI](#)

Large Language Models Meet **Stance Detection**: A Survey of Tasks, Methods, Applications, Challenges and Future Directions

Authors: [Lata Pangtey](#), [Anukriti Bhatnagar](#), [Shubhi Bansal](#), [Shahid Shafi Dar](#), [Nagendra Kumar](#)

Abstract: **Stance**... [▽ More](#)

Submitted 19 January, 2026; v1 submitted 13 May, 2025; originally announced May 2025.

148. [arXiv:2505.05459](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.SI](#)

UKElectionNarratives: A Dataset of Misleading Narratives Surrounding Recent UK General Elections

Authors: [Fatima Haouari](#), [Carolina Scarton](#), [Nicolò Faggiani](#), [Nikolaos Nikolaidis](#), [Bonka Kotseva](#), [Ibrahim Abu Farha](#), [Jens Linge](#), [Kalina Bontcheva](#)

Abstract: ...misleading narratives that circulated during recent elections in Europe. Based on this taxonomy, we construct and analyse UKElectionNarratives: the first dataset of human-**annotated** misleading narratives which circulated during the UK General Elections in 2019 and 2024. We also...

[▽ More](#)

Submitted 8 May, 2025; **originally announced** May 2025.

Comments: This work was accepted at the International AAAI Conference on Web and **Social Media** (ICWSM 2025)

149. [arXiv:2505.05148](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

A **Benchmark** Dataset and a Framework for Urdu Multimodal Named Entity Recognition

Authors: [Hussain Ahmad](#), [Qingyang Zeng](#), [Jing Wan](#)

Abstract: The emergence of multimodal content, particularly text and images on **social**... [▽ More](#)

Submitted 8 May, 2025; **originally announced** May 2025.

Comments: 16 pages, 5 figures. Preprint

150. [arXiv:2504.21026](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#) [cs.SI](#)

Creating and Evaluating Code-Mixed Nepali-English and Telugu-English Datasets for Abusive Language Detection Using Traditional and Deep Learning Models

Authors: [Manish Pandey](#), [Nageshwar Prasad Yadav](#), [Mokshada Adduru](#), [Sawan Rai](#)

Abstract: With the growing presence of multilingual users on **social**... [▽ More](#)

Submitted 23 April, 2025; **originally announced** April 2025.

151. [arXiv:2504.15987](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.CY](#)

Few-shot Hate Speech Detection Based on the MindSpore Framework

Authors: [Zhenkai Qin](#), [Dongze Wu](#), [Yuxin Liu](#), [Guifang Yang](#)

Abstract: The proliferation of hate speech on **social**... [▽ More](#)

Submitted 22 April, 2025; **originally announced** April 2025.

152. [arXiv:2504.15548](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.CY](#)

LLM-based Semantic Augmentation for Harmful Content Detection

Authors: [Elyas Meguellati](#), [Assaad Zeghina](#), [Shazia Sadiq](#), [Gianluca Demartini](#)

Abstract: Recent advances in **large**... [▽ More](#)

Submitted 21 April, 2025; **originally announced** April 2025.

153. [arXiv:2504.13231](#) [[pdf](#), [ps](#), [other](#)] [cs.CV](#) [cs.AI](#)

WildFireCan-MMD: A Multimodal Dataset for Classification of **User-Generated Content** During Wildfires in Canada

Authors: [Braeden Sherritt](#), [Isar Nejadgholi](#), [Efstratios Aivaliotis](#), [Khaled Mslmani](#), [Marzieh Amini](#)

Abstract: Rapid information access is vital during wildfires, yet traditional data sources are slow and costly. **Social**... [▽ More](#)

Submitted 10 November, 2025; **v1** submitted 17 April, 2025; **originally announced** April 2025.

154. [arXiv:2504.12312](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Socrates or Smartypants: Testing Logic Reasoning Capabilities of **Large Language Models** with Logic Programming-based Test Oracles

Authors: [Zihao Xu](#), [Junchen Ding](#), [Yiling Lou](#), [Kun Zhang](#), [Dong Gong](#), [Yuekang Li](#)

Abstract: [Large...](#) ▾ More

Submitted 18 November, 2025; v1 submitted 9 April, 2025; originally announced April 2025.

155. [arXiv:2504.09958](#) [[pdf](#), [ps](#), [other](#)] cs.CL

C-MTCSD: A Chinese Multi-Turn Conversational [Stance Detection](#) Dataset

Authors: [Fuqiang Niu](#), [Yi Yang](#), [Xianghua Fu](#), [Genan Dai](#), [Bowen Zhang](#)

Abstract: [Stance...](#) ▾ More

Submitted 18 April, 2025; v1 submitted 14 April, 2025; originally announced April 2025.

Comments: WWW2025

156. [arXiv:2504.06166](#) [[pdf](#), [other](#)] cs.CL

Assessing how hyperparameters impact [Large Language Models'](#) sarcasm detection performance

Authors: [Montgomery Gole](#), [Andriy Miranskyy](#)

Abstract: ...impact sarcasm detection in OpenAI's GPT, and Meta's Llama-2 models, given their strong natural language understanding, and popularity. We evaluate fine-tuned and [zero-](#)... ▾ More

Submitted 15 April, 2025; v1 submitted 8 April, 2025; originally announced April 2025.

Comments: arXiv admin note: substantial text overlap with arXiv:2312.04642

157. [arXiv:2504.04295](#) [[pdf](#), [other](#)] cs.CL cs.CE

Dynamic Hedging Strategies in Derivatives Markets with LLM-Driven Sentiment and News Analytics

Authors: [Jie Yang](#), [Yiqiu Tang](#), [Yongjie Li](#), [Lihua Zhang](#), [Haoran Zhang](#)

Abstract: ...effective risk management in derivatives markets, where volatility and market sentiment can greatly impact performance. This paper introduces a novel framework that leverages [large](#)... ▾ More

Submitted 5 April, 2025; originally announced April 2025.

Comments: Accepted by IJCNN 2025

158. [arXiv:2503.21679](#) [[pdf](#), [ps](#), [other](#)] cs.CL cs.CY

JiraiBench: A Bilingual [Benchmark](#) for Evaluating [Large Language Models'](#) Detection of Human Self-Destructive Behavior Content in Jirai Community

Authors: [Yunze Xiao](#), [Tingyu He](#), [Lionel Z. Wang](#), [Yiming Ma](#), [Xingyu Song](#), [Xiaohang Xu](#), [Mona Diab](#), [Irene Li](#), [Ka Chung Ng](#)

Abstract: This paper introduces JiraiBench, the first bilingual [benchmark](#) for evaluating... ▾ More

Submitted 25 January, 2026; v1 submitted 27 March, 2025; originally announced March 2025.

Comments: 20 pages, 1 figures

159. [arXiv:2503.18253](#) [[pdf](#), [ps](#), [other](#)] cs.CL

Enhancing Multi-Label Emotion Analysis and Corresponding Intensities for Ethiopian Languages

Authors: [Tadesse Destaw Belay](#), [Dawit Ketema Gete](#), [Abinew Ali Ayele](#), [Olga Kolesnikova](#), [Iqra Ameer](#), [Grigori Sidorov](#), [Seid Muhie Yimam](#)

Abstract: ...and integrating emotion-understanding models are essential for a wide range of human-computer interaction tasks, including customer feedback analysis, marketing research, and

[social...](#) ▾ [More](#)

Submitted 19 March, 2026; v1 submitted 23 March, 2025; originally announced March 2025.

Comments: LREC 2026

160. [arXiv:2503.18117](#) [[pdf](#), [other](#)] [cs.CL](#)

Detection of Somali-written Fake News and Toxic Messages on the **Social Media** Using Transformer-based Language Models

Authors: [Muhidin A. Mohamed](#), [Shuab D. Ahmed](#), [Yahye A. Isse](#), [Hanad M. Mohamed](#), [Fuad M. Hassan](#), [Houssein A. Assowe](#)

Abstract: The fact that everyone with a [social...](#) ▾ [More](#)

Submitted 23 March, 2025; originally announced March 2025.

161. [arXiv:2503.15169](#) [[pdf](#)] [cs.CL](#) [cs.AI](#)

Benchmarking Open-Source **Large Language Models** on Healthcare Text Classification Tasks

Authors: [Yuting Guo](#), [Abeed Sarker](#)

Abstract: The application of [large...](#) ▾ [More](#)

Submitted 8 May, 2025; v1 submitted 19 March, 2025; originally announced March 2025.

Comments: 5 pages

162. [arXiv:2503.10095](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Cognitive-Mental-LLM: Evaluating Reasoning in **Large Language Models** for Mental Health Prediction via Online Text

Authors: [Avinash Patil](#), [Amardeep Kour Gedhu](#)

Abstract: [Large...](#) ▾ [More](#)

Submitted 7 January, 2026; v1 submitted 13 March, 2025; originally announced March 2025.

Comments: 8 pages, 4 Figures, 3 tables

163. [arXiv:2503.05142](#) [[pdf](#), [other](#)] [cs.CL](#)

RocketEval: Efficient Automated LLM Evaluation via Grading Checklist

Authors: [Tianjun Wei](#), [Wei Wen](#), [Ruizhi Qiao](#), [Xing Sun](#), [Jianghong Ma](#)

Abstract: Evaluating [large...](#) ▾ [More](#)

Submitted 6 March, 2025; originally announced March 2025.

Comments: Accepted by ICLR 2025: <https://openreview.net/forum?id=zJjzNj6QUe>

164. [arXiv:2502.19954](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Collaborative **Stance Detection** via Small-Large Language Model Consistency Verification

Authors: [Yu Yan](#), [Sheng Sun](#), [Zixiang Tang](#), [Teli Liu](#), [Min Liu](#)

Abstract: [Stance...](#) ▾ [More](#)

Submitted 22 August, 2025; v1 submitted 27 February, 2025; originally announced February 2025.

165. [arXiv:2502.16503](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.LG](#)

FanChuan: A Multilingual and Graph-Structured **Benchmark** For Parody Detection and Analysis

Authors: [Yilun Zheng](#), [Sha Li](#), [Fangkun Wu](#), [Yang Ziyi](#), [Lin Hongchao](#), [Zhichao Hu](#), [Cai Xinjun](#), [Ziming Wang](#), [Jinxuan Chen](#), [Sitao Luan](#), [Jiahao Xu](#), [Lihui Chen](#)

Abstract: Parody is an emerging phenomenon on **social**... [▽ More](#)

Submitted 23 February, 2025; **originally announced** February 2025.

166. [arXiv:2502.16101](#) [[pdf](#), [ps](#), [other](#)] [cs.AI](#) [cs.IR](#)

Worse than **Zero-shot**? A Fact-Checking Dataset for Evaluating the Robustness of RAG Against Misleading Retrievals

Authors: [Linda Zeng](#), [Rithwik Gupta](#), [Divij Motwani](#), [Yi Zhang](#), [Diji Yang](#)

Abstract: Retrieval-augmented generation (RAG) has shown impressive capabilities in mitigating hallucinations in **large**... [▽ More](#)

Submitted 16 January, 2026; **v1** submitted 22 February, 2025; **originally announced** February 2025.

Comments: Advances in Neural Information Processing Systems (NeurIPS 2025)

167. [arXiv:2502.14083](#) [[pdf](#), [other](#)] [cs.CL](#)

Are Rules Meant to be Broken? Understanding Multilingual Moral Reasoning as a Computational Pipeline with UniMoral

Authors: [Shivani Kumar](#), [David Jurgens](#)

Abstract: ...discordant datasets and tasks that examine isolated aspects of moral reasoning. We bridge this gap with UniMoral, a unified dataset integrating psychologically grounded and **social**... [▽ More](#)

Submitted 19 February, 2025; **originally announced** February 2025.

Comments: 21 pages, 10 figures, 8 tables

168. [arXiv:2502.06150](#) [[pdf](#), [other](#)] [cs.CL](#)

Scaling Public Health Text **Annotation: Zero-Shot** Learning vs. Crowdsourcing for Improved Efficiency and Labeling Accuracy

Authors: [Kamyar Kazari](#), [Yong Chen](#), [Zahra Shakeri](#)

Abstract: Public health researchers are increasingly interested in using **social**... [▽ More](#)

Submitted 9 February, 2025; **originally announced** February 2025.

Comments: 4 pages, 1 figure

169. [arXiv:2502.04346](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [doi](#) [10.1109/TCSS.2025.3623021](#)

Multi-Lingual Cyber Threat Detection in Tweets/X Using ML, DL, and LLM: A **Comparative** Analysis

Authors: [Saydul Akbar Murad](#), [Ashim Dahal](#), [Nick Rahimi](#)

Abstract: Cyber threat detection has become an important area of focus in today's digital age due to the growing spread of fake information and harmful content on **social**... [▽ More](#)

Submitted 3 February, 2025; **originally announced** February 2025.

Journal ref: IEEE Trans. Comput. Soc. Syst., Early Access, pp. 1-15, Nov. 2025

170. [arXiv:2502.01658](#) [[pdf](#)] [cs.CL](#) [cs.CY](#) [cs.SI](#) [doi](#) [10.2196/63631](#)

Large Language Models' Accuracy in Emulating Human Experts' Evaluation of Public Sentiments about Heated Tobacco Products on **Social Media**

Authors: [Kwanho Kim](#), [Soojong Kim](#)

Abstract: **Sentiment**... [▽ More](#)

171. [arXiv:2502.00414](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Social media polarization during conflict: Insights from an ideological stance dataset on Israel-Palestine **Reddit** comments

Authors: [Hasin Jawad Ali](#), [Ajwad Abrar](#), [S. M. Hozafa Hossain](#), [M. Firoz Mridha](#)

Abstract: In politically sensitive scenarios like wars, **social**... [▽ More](#)

Submitted 15 April, 2026; v1 submitted 1 February, 2025; originally announced February 2025.

172. [arXiv:2501.16836](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

Misspellings in Natural Language Processing: A survey

Authors: [Gianluca Sperduti](#), [Alejandro Moreo](#)

Abstract: ...in natural language processing (NLP). While often unintentional, misspellings have become ubiquitous in digital communication, especially with the proliferation of Web 2.0, **user**... [▽ More](#)

Submitted 24 October, 2025; v1 submitted 28 January, 2025; originally announced January 2025.

173. [arXiv:2501.14037](#) [[pdf](#)] [cs.CL](#)

Emotions, Context, and Substance Use in Adolescents: A **Large Language Model** Analysis of **Reddit** Posts

Authors: [Jianfeng Zhu](#), [Hailong Jiang](#), [Yulan Wang](#), [Karin G. Coifman](#), [Ruoming Jin](#), [Deric R. Kenne](#)

Abstract: ...contextual factors driving these behaviors remain poorly understood. This study analyzed 23000 substance-use related posts and an equal number of non-substance posts from **Reddit's** r/teenagers community (2018-2022). Posts were... [▽ More](#)

Submitted 14 November, 2025; v1 submitted 23 January, 2025; originally announced January 2025.

Comments: 19 pages, 5 figures

174. [arXiv:2501.13802](#) [[pdf](#), [other](#)] [cs.CY](#)

Enhancing LLMs for Governance with Human Oversight: Evaluating and Aligning LLMs on Expert Classification of Climate Misinformation for Detecting False or Misleading Claims about Climate Change

Authors: [Mowafak Allaham](#), [Ayse D. Lokmanoglu](#), [P. Sol Hart](#), [Erik C. Nisbet](#)

Abstract: Climate misinformation is a problem that has the potential to be substantially aggravated by the development of **Large**... [▽ More](#)

Submitted 9 March, 2025; v1 submitted 23 January, 2025; originally announced January 2025.

Comments: International Workshop on AI Governance: Alignment, Morality and Law (AIGOV) 2025. AAAI Conference on Artificial Intelligence

175. [arXiv:2501.11770](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.CY](#) [cs.SI](#)

The Value of Nothing: Multimodal Extraction of Human Values Expressed by TikTok Influencers

Authors: [Alina Starovolsky-Shitrit](#), [Alon Neduva](#), [Naama Appel Doron](#), [Itamar Gafni](#), [Ella Daniel](#), [Oren Tsur](#)

Abstract: ...paper we extract implicit values from TikTok movies uploaded by online influencers targeting children and adolescents. We curated a dataset of hundreds of TikTok movies and **annotated** them according to the well established Schwartz Theory of Personal Values. We then

experimented with an array of language models, investigating their utility in value identifica...

[▽ More](#)

Submitted 26 March, 2026; v1 submitted 20 January, 2025; originally announced January 2025.

176. [arXiv:2501.08413](#) [[pdf](#), [ps](#), [other](#)] cs.CL

Labeling Free-text Data using Language Model Ensembles

Authors: [Jiaxing Qiu](#), [Dongliang Guo](#), [Natalie Papini](#), [Noelle Peace](#), [Hannah F. Fitterman-Harris](#), [Cheri A. Levinson](#), [Tom Hartvigsen](#), [Teague R. Henry](#)

Abstract: ...measures may not capture. Labeling curated topics of research interest in free-text data by multiple trained human coders is typically labor-intensive and time-consuming. Though **large**...

[▽ More](#)

Submitted 25 September, 2025; v1 submitted 14 January, 2025; originally announced January 2025.

177. [arXiv:2501.01256](#) [[pdf](#), [other](#)] cs.CL cs.LG

Digital Guardians: Can GPT-4, Perspective API, and Moderation API reliably detect hate speech in reader comments of German online newspapers?

Authors: [Manuel Weber](#), [Moritz Huber](#), [Maximilian Auch](#), [Alexander Döschl](#), [Max-Emanuel Keller](#), [Peter Mandl](#)

Abstract: In recent years, toxic content and hate speech have become widespread phenomena on the internet. Moderators of online newspapers and **forums** are now required, partly due to legal regulations, to carefully review and, if necessary, delete reader comments. This is a labor-intensive process. Some providers of... [▽ More](#)

Submitted 2 January, 2025; originally announced January 2025.

ACM Class: I.2.7

178. [arXiv:2412.19928](#) [[pdf](#)] cs.CL cs.SI

Assessing **Text Classification** Methods for Cyberbullying Detection on **Social Media** Platforms

Authors: [Adamu Gaston Philipo](#), [Doreen Sebastian Sarwatt](#), [Jianguo Ding](#), [Mahmoud Daneshmand](#), [Huansheng Ning](#)

Abstract: Cyberbullying significantly contributes to mental health issues in communities by negatively impacting the psychology of victims. It is a prevalent problem on **social**... [▽ More](#)

Submitted 27 December, 2024; originally announced December 2024.

Comments: 15 pages, 10 figures, 7 tables

179. [arXiv:2412.14008](#) [[pdf](#)] cs.CL

FarExStance: Explainable **Stance Detection** for Farsi

Authors: [Majid Zarharan](#), [Maryam Hashemi](#), [Malika Behroozrazegh](#), [Sauleh Eetemadi](#), [Mohammad Taher Pilehvar](#), [Jennifer Foster](#)

Abstract: We introduce FarExStance, a new dataset for explainable **stance**... [▽ More](#)

Submitted 18 December, 2024; originally announced December 2024.

Comments: Accepted in COLING 2025

180. [arXiv:2412.06951](#) [[pdf](#), [other](#)] cs.CL cs.LG

Analysing Public Transport User Sentiment on **Low Resource** Multilingual Data

Authors: [Rozina L. Myoya](#), [Vukosi Marivate](#), [Idris Abdulmumin](#)

Abstract: Public transport systems in many Sub-Saharan countries often receive less attention **compared** to other sectors, underscoring the need for innovative solutions to improve the Quality of Service (QoS) and overall user experience. This study explored commuter opinion mining to understand sentiments toward existing public transport systems in Kenya, Tanzania, and... [▽ More](#)
Submitted 9 December, 2024; **originally announced** December 2024.

181. [arXiv:2412.04726](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

BESSTIE: A **Benchmark** for Sentiment and Sarcasm Classification for Varieties of English

Authors: [Dipankar Srirag](#), [Aditya Joshi](#), [Jordan Painter](#), [Diptesh Kanojia](#)

Abstract: Despite **large**... [▽ More](#)

Submitted 17 June, 2025; **v1** submitted 5 December, 2024; **originally announced** December 2024.

Comments: Findings of ACL: ACL 2025

182. [arXiv:2412.02271](#) [[pdf](#)] [cs.CL](#)

The MediaSpin Dataset: Post-Publication News Headline Edits **Annotated** for Media Bias

Authors: [Preetika Verma](#), [Kokil Jaidka](#)

Abstract: ...how major news outlets modify headlines after publication, and MediaSpin-in-the-Wild, a complementary dataset linking these revised headlines to their downstream engagement on **social**... [▽ More](#)

Submitted 14 May, 2026; **v1** submitted 3 December, 2024; **originally announced** December 2024.

Comments: 8 pages, 3 figures, 8 tables Accepted at AAAI ICWSM 2026 We updated the paper title from "MediaSpin: Exploring Media Bias Through Fine-Grained Analysis of News Headlines " to "The MediaSpin Dataset: Post-Publication News Headline Edits Annotated for Media Bias"

183. [arXiv:2411.19832](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#)

Sensitive Content Classification in **Social Media**: A Holistic Resource and Evaluation

Authors: [Dimosthenis Antypas](#), [Indira Sen](#), [Carla Perez-Almendros](#), [Jose Camacho-Collados](#), [Francesco Barbieri](#)

Abstract: ...on toxic language, leaving gaps in detecting other sensitive categories such as substance abuse or self-harm. In this paper, we put forward a unified dataset tailored for **social**... [▽ More](#)

Submitted 24 June, 2025; **v1** submitted 29 November, 2024; **originally announced** November 2024.

Comments: Accepted at the 9th Workshop on Online Abuse and Harms (WOAH)

ACM Class: I.2.7

184. [arXiv:2411.09612](#) [[pdf](#)] [cs.CL](#) [cs.LG](#)

The Moral Foundations Weibo Corpus

Authors: [Renjie Cao](#), [Miaoyan Hu](#), [Jiahan Wei](#), [Baha Ilnaini](#)

Abstract: Moral sentiments expressed in natural language significantly influence both online and offline environments, shaping behavioral styles and interaction patterns, including **social**... [▽ More](#)

Submitted 14 November, 2024; **originally announced** November 2024.

Report number: UNI-2024-157

185. [arXiv:2411.02666](#) [[pdf](#), [other](#)] [cs.LG](#) [cs.AI](#) [cs.SI](#)

From **Twitter** to Reasoner: Understand Mobility Travel Modes and Sentiment Using **Large Language Models**

Authors: [Kangrui Ruan](#), [Xinyang Wang](#), [Xuan Di](#)

Abstract: **Social...** [▽ More](#)

Submitted 4 November, 2024; **originally announced** November 2024.

Comments: 6 pages; Accepted by ITSC 2024

186. [arXiv:2410.15016](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [cs.CY](#) [cs.IR](#) [cs.LG](#) [cs.SI](#)

Transit Pulse: Utilizing **Social Media** as a Source for Customer Feedback and Information Extraction with **Large Language Model**

Authors: [Jiahao Wang](#), [Amer Shalaby](#)

Abstract: ...makes manual analysis impractical, and standard NLP techniques like Term Frequency-Inverse Document Frequency (TF-IDF) fall short in nuanced interpretation. Traditional **sentiment...** [▽ More](#)

Submitted 19 October, 2024; **originally announced** October 2024.

Comments: 17 pages, 21 figures

187. [arXiv:2410.13604](#) [[pdf](#), [other](#)] [cs.IR](#) [cs.AI](#) [cs.CL](#)

Large Language Models as Narrative-Driven Recommenders

Authors: [Lukas Eberhard](#), [Thorsten Ruprechter](#), [Denis Helic](#)

Abstract: ...personalized suggestions for user requests expressed in free-form text such as "I want to watch a thriller with a mind-bending story, like Shutter Island." Although **large...** [▽ More](#)

Submitted 17 October, 2024; **originally announced** October 2024.

Comments: Under review; 19 pages

188. [arXiv:2410.08900](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

A **Benchmark** for Cross-Domain Argumentative Stance Classification on **Social Media**

Authors: [Jiaqing Yuan](#), [Ruijie Xi](#), [Munindar P. Singh](#)

Abstract: ...role in identifying authors' viewpoints on specific topics. However, generating diverse pairs of argumentative sentences across various domains is challenging. Existing **benchmarks** often come from a single domain or focus on a limited set of topics. Additionally, manual... [▽ More](#)

Submitted 15 November, 2024; **v1** submitted 11 October, 2024; **originally announced** October 2024.

Comments: Accepted by AAAI ICWSM 2025

189. [arXiv:2410.06338](#) [[pdf](#), [other](#)] [cs.CL](#)

Are **Large Language Models** State-of-the-art Quality Estimators for Machine Translation of **User-generated Content**?

Authors: [Shenbin Qian](#), [Constantin Orăsan](#), [Diptesh Kanojia](#), [Félix do Carmo](#)

Abstract: This paper investigates whether **large...** [▽ More](#)

Submitted 8 October, 2024; **originally announced** October 2024.

190. [arXiv:2410.03908](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#) [doi](#) [10.18653/v1/2024.emnlp-main.931](#)

Still Not Quite There! Evaluating **Large Language Models** for Comorbid Mental Health Diagnosis

Authors: [Amey Hengle](#), [Atharva Kulkarni](#), [Shantanu Patankar](#), [Madhumitha Chandrasekaran](#), [Sneha D'Silva](#), [Jemima Jacob](#), [Rashmi Gupta](#)

Abstract: In this study, we introduce ANGST, a novel, first-of-its kind **benchmark** for depression-anxiety comorbidity classification from... [▽ More](#)

Submitted 4 October, 2024; **originally announced** October 2024.

Comments: 24 Pages

191. [arXiv:2410.02084](#) [[pdf](#), [ps](#), [other](#)] [cs.SD](#) [eess.AS](#)

Generating Symbolic Music from Natural Language Prompts using an LLM-Enhanced Dataset

Authors: [Weihan Xu](#), [Julian McAuley](#), [Taylor Berg-Kirkpatrick](#), [Shlomo Dubnov](#), [Hao-Wen Dong](#)

Abstract: ...with extensive metadata and captions. In this work, we present MetaScore, a new dataset consisting of 963K musical scores paired with rich metadata, including free-form user-**annotated** tags, collected from an online music... [▽ More](#)

Submitted 16 June, 2025; **v1** submitted 2 October, 2024; **originally announced** October 2024.

Comments: Accepted at ISMIR 2025

192. [arXiv:2409.19492](#) [[pdf](#), [ps](#), [other](#)] [cs.CL](#) [cs.AI](#)

MedHalu: Hallucinations in Responses to Healthcare Queries by Large Language Models

Authors: [Vibhor Agarwal](#), [Yiqiao Jin](#), [Mohit Chandra](#), [Munmun De Choudhury](#), [Srijan Kumar](#), [Nishanth Sastry](#)

Abstract: **Large**... [▽ More](#)

Submitted 22 November, 2025; **v1** submitted 28 September, 2024; **originally announced** September 2024.

Comments: Accepted at ICWSM2026. <https://netsys.surrey.ac.uk/datasets/medhalu/>

193. [arXiv:2409.16807](#) [[pdf](#), [other](#)] [cs.CL](#)

A Few Hypocrites: **Few-Shot** Learning and Subtype Definitions for Detecting Hypocrisy Accusations in Online Climate Change Debates

Authors: [Paulina Garcia Corral](#), [Avishai Green](#), [Hendrik Meyer](#), [Anke Stoll](#), [Xiaoyue Yan](#), [Myrthe Reuver](#)

Abstract: ...detection as an independent task in NLP, and identify different relevant subtypes of hypocrisy accusations. Our Climate Hypocrisy Accusation Corpus (CHAC) consists of 420 **Reddit** climate debate comments, expert-... [▽ More](#)

Submitted 25 September, 2024; **originally announced** September 2024.

Comments: cite the public version, published at CPSS 2024 @ KONVENS

194. [arXiv:2409.09741](#) [[pdf](#), [other](#)] [cs.CL](#) [cs.AI](#)

Benchmarking LLMs in Political Content Text-**Annotation**: Proof-of-Concept with Toxicity and Incivility Data

Authors: [Bastián González-Bustamante](#)

Abstract: This article **benchmarked** the ability of OpenAI's GPTs and a number of open-source LLMs to perform... [▽ More](#)

Submitted 15 September, 2024; **originally announced** September 2024.

Comments: Paper prepared for delivery at the 8th Monash-Warwick-Zurich Text-as-Data Workshop, September 16-17, 2024: 11 pages, 3 tables, 3 figures

MSC Class: 68T50 (Primary) 91F10; 91F20 (Secondary)

195. [arXiv:2409.09467](#) [pdf, other] cs.CL

Keeping Humans in the Loop: Human-Centered Automated Annotation with Generative AI

Authors: [Nicholas Pangakis](#), [Samuel Wolken](#)

Abstract: Automated text annotation is a compelling use case for generative... [▽ More](#)

Submitted 21 September, 2024; v1 submitted 14 September, 2024; originally announced September 2024.

Comments: Accepted at Proceedings of the International AAAI Conference on Web and Social Media. Vol. 19. 2025

196. [arXiv:2408.02044](#) [pdf, other] cs.CL cs.AI

Fine-tuning multilingual language models in Twitter/X sentiment analysis: a study on Eastern-European V4 languages

Authors: [Tomáš Filip](#), [Martin Pavlíček](#), [Petr Sosík](#)

Abstract: The aspect-based sentiment... [▽ More](#)

Submitted 4 August, 2024; originally announced August 2024.

Comments: 18 pages, 4 figures

197. [arXiv:2407.12141](#) [pdf, other] cs.CL cs.AI

Predicting Emotion Intensity in Polish Political Texts: Comparing Supervised Models and Large Language Models in a Resource-Poor Language

Authors: [Hubert Plisiecki](#), [Piotr Koc](#), [Maria Flakus](#), [Artur Pokropek](#)

Abstract: This study explores the use of large... [▽ More](#)

Submitted 16 July, 2024; originally announced July 2024.

Comments: The Appendix is located at the very bottom of the manuscript

198. [arXiv:2407.07258](#) [pdf, other] cs.CL cs.LG

Identification of emotions on Twitter during the 2022 electoral process in Colombia

Authors: [Juan Jose Iguaran Fernandez](#), [Juan Manuel Perez](#), [German Rosati](#)

Abstract: The study of Twitter as a means for analyzing social phenomena has gained interest in recent years due to the availability of large amounts of data in a relatively spontaneous environment. Within opinion-mining tasks, emotion detection is specially relevant, as it allows for the identification of people's subjective responses to different social events i... [▽ More](#)

Submitted 9 July, 2024; originally announced July 2024.

199. [arXiv:2407.00369](#) [pdf, other] cs.CL

How to Train Your Fact Verifier: Knowledge Transfer with Multimodal Open Models

Authors: [Jaeyoung Lee](#), [Ximing Lu](#), [Jack Hessel](#), [Faeze Brahman](#), [Youngjae Yu](#), [Yonatan Bisk](#), [Yejin Choi](#), [Saadia Gabriel](#)

Abstract: Given the growing influx of misinformation across news and social... [▽ More](#)

Submitted 29 June, 2024; originally announced July 2024.

200. [arXiv:2407.00191](#) [pdf, other] cs.CL

MetaKP: On-Demand Keyphrase Generation

Authors: [Di Wu](#), [Xiaoxian Shen](#), [Kai-Wei Chang](#)

Abstract: ...on-demand keyphrase generation, a novel paradigm that requires keyphrases that conform to specific high-level goals or intents. For this task, we present MetaKP, a large-scale **benchmark** comprising four datasets, 7500 documents, and 3760 goals across news and biomedical domains with human-... [▽ More](#)

Submitted 4 October, 2024; **v1** submitted 28 June, 2024; **originally announced** July 2024.

Comments: EMNLP 2024 (Findings)

1

2

Next

We gratefully acknowledge support from our **major funders**, **member institutions**, and all contributors.

[About](#) · [Help](#) · [Contact](#) · [Subscribe](#) · [Copyright](#) · [Privacy](#) · [Accessibility](#) · [Operational Status](#)

Major funding support from

